自然語言處理與應用
Natural Language Processing and Applications

Transformer

Instructor: 林英嘉 (Ying-Jia Lin)
2026/03/23

GitHub

Slido
NLP0323
(Link)

Outline

- Recaps [10 min]
- Transformers [90 min]
- Homework 2 [30 min]
- Quiz [15 min]

深度學習常見模型架構

CNN
(著重局部資訊)

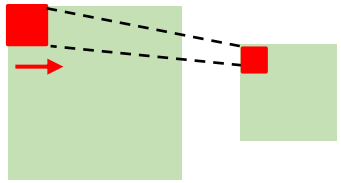

RNN
(著重記憶力)

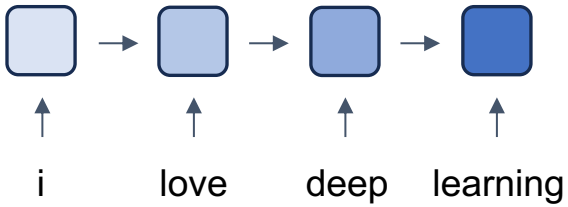

Transformer
(著重全局資訊與記憶力)

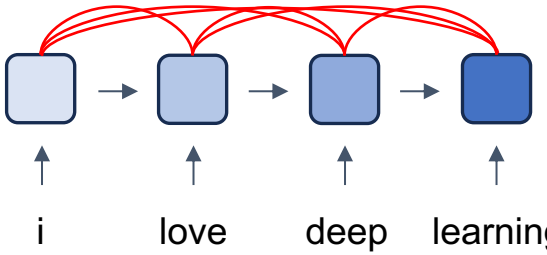

[Recap] Recurrent Neural Networks (RNN)

- Recurrent Neural Networks (RNNs) are a type of artificial neural network designed to handle sequential data by capturing temporal dependencies.
  - RNN 的“遞迴”是指它在每一個時間步中重複使用相同的參數（同一組權重），並把先前的隱藏狀態傳到下一個時間步

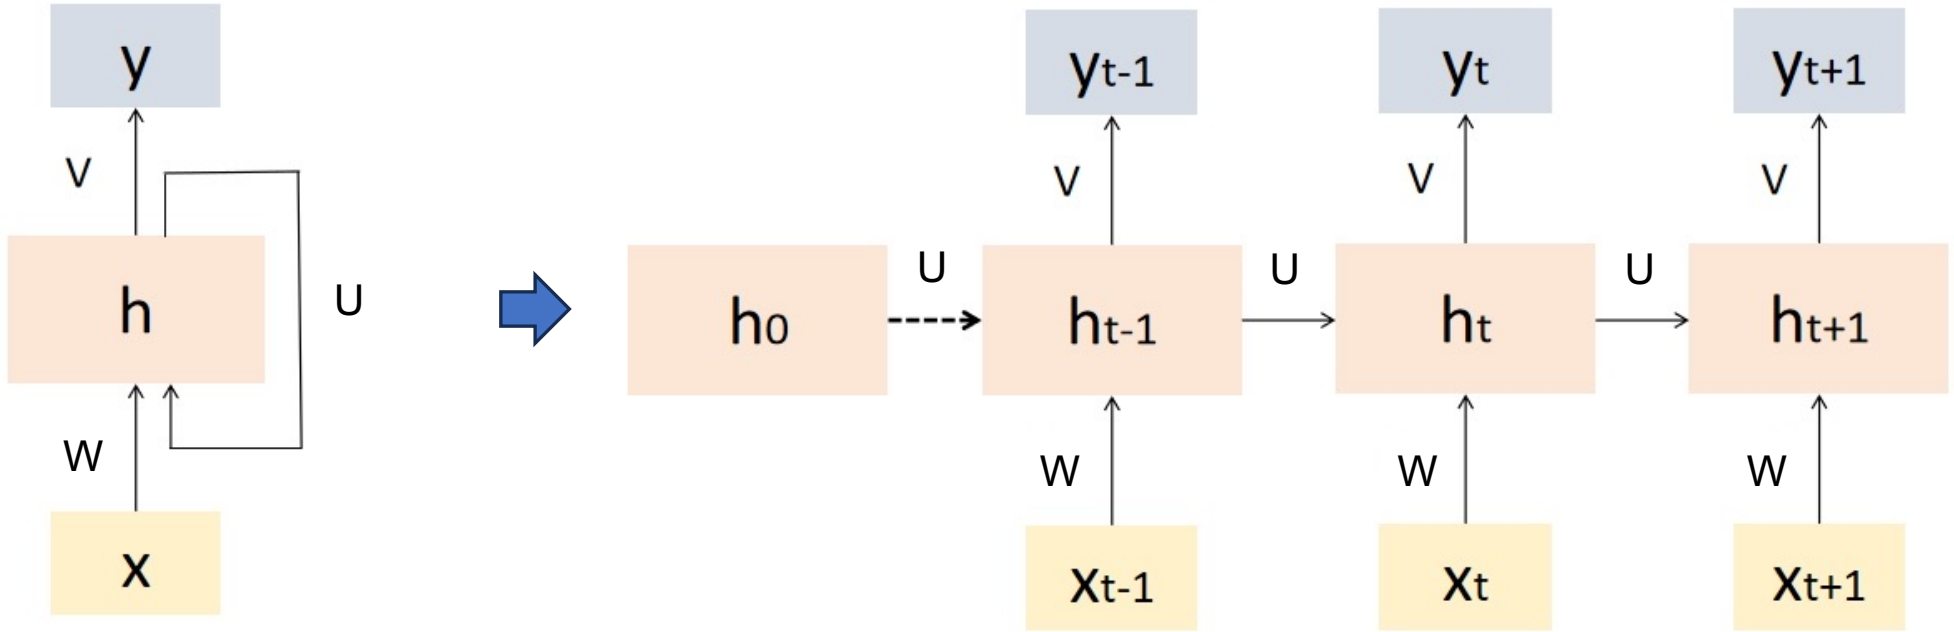

[Recap] Text Generation with RNN

*本頁的RNN不包含輸出層V

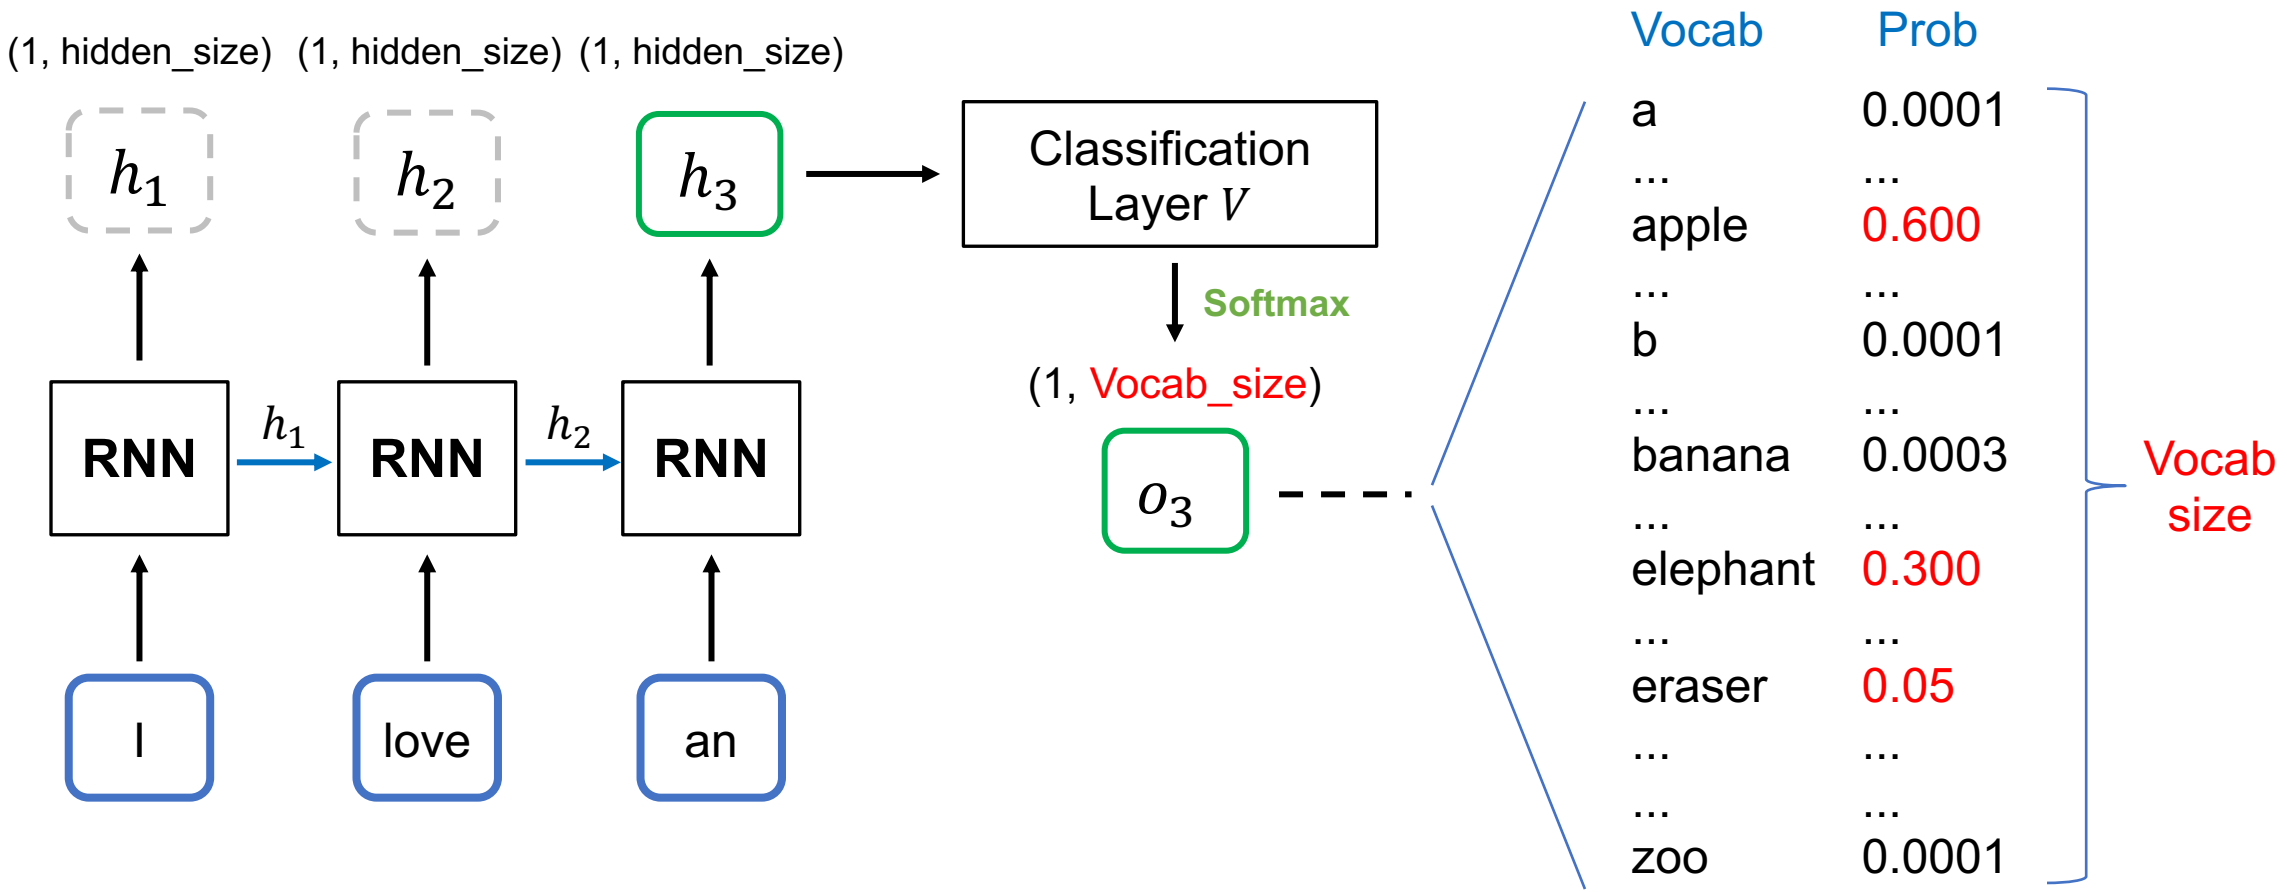

NLP

Word embeddings

[Recap] RNN 的記憶力意義

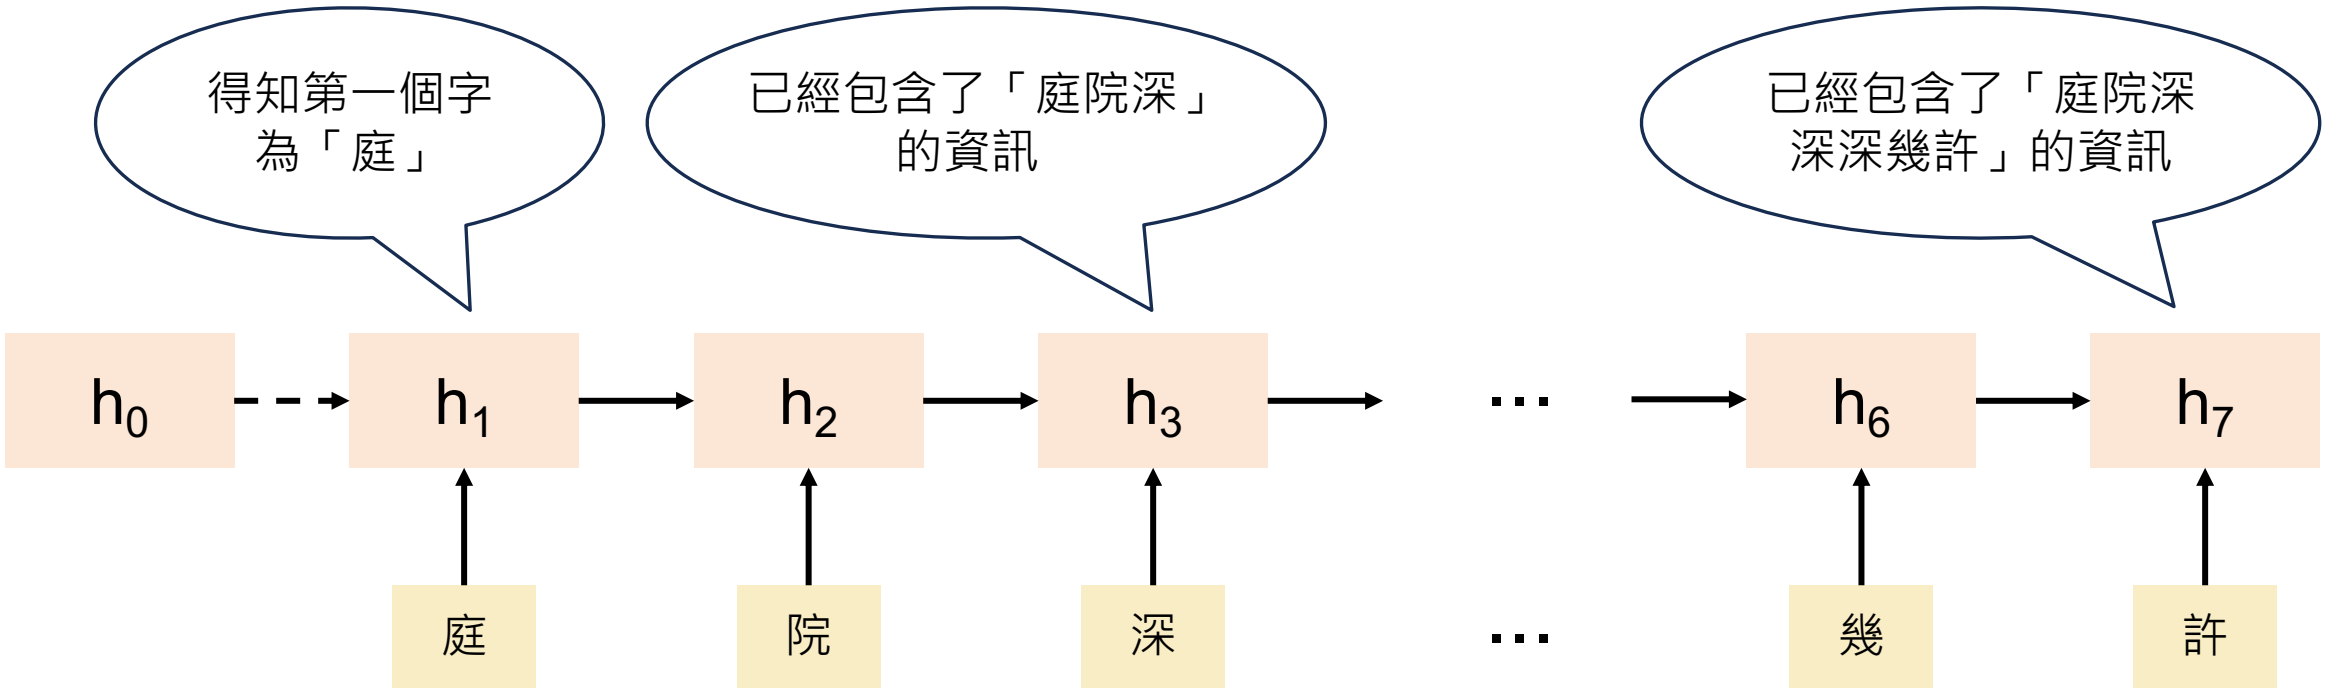

為什麼我們需要序列模型?

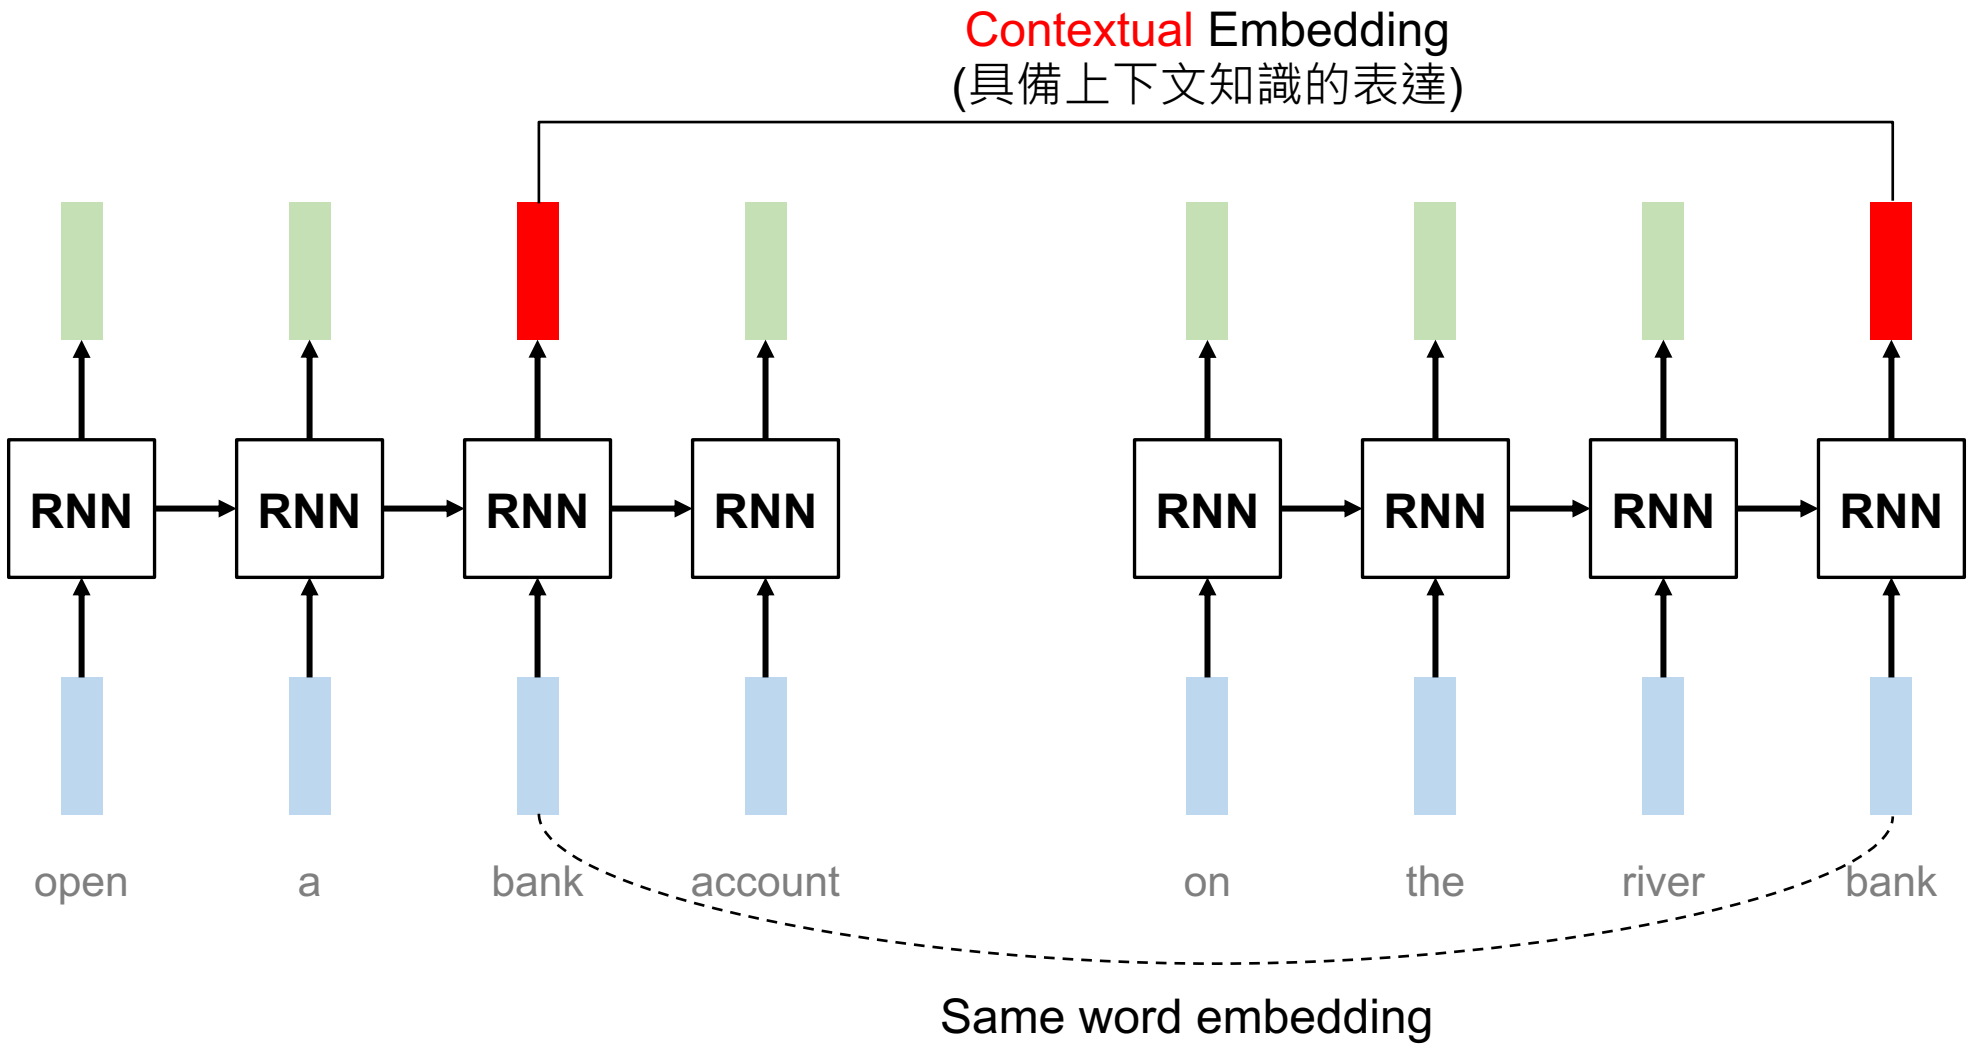

Encoder and Decoder

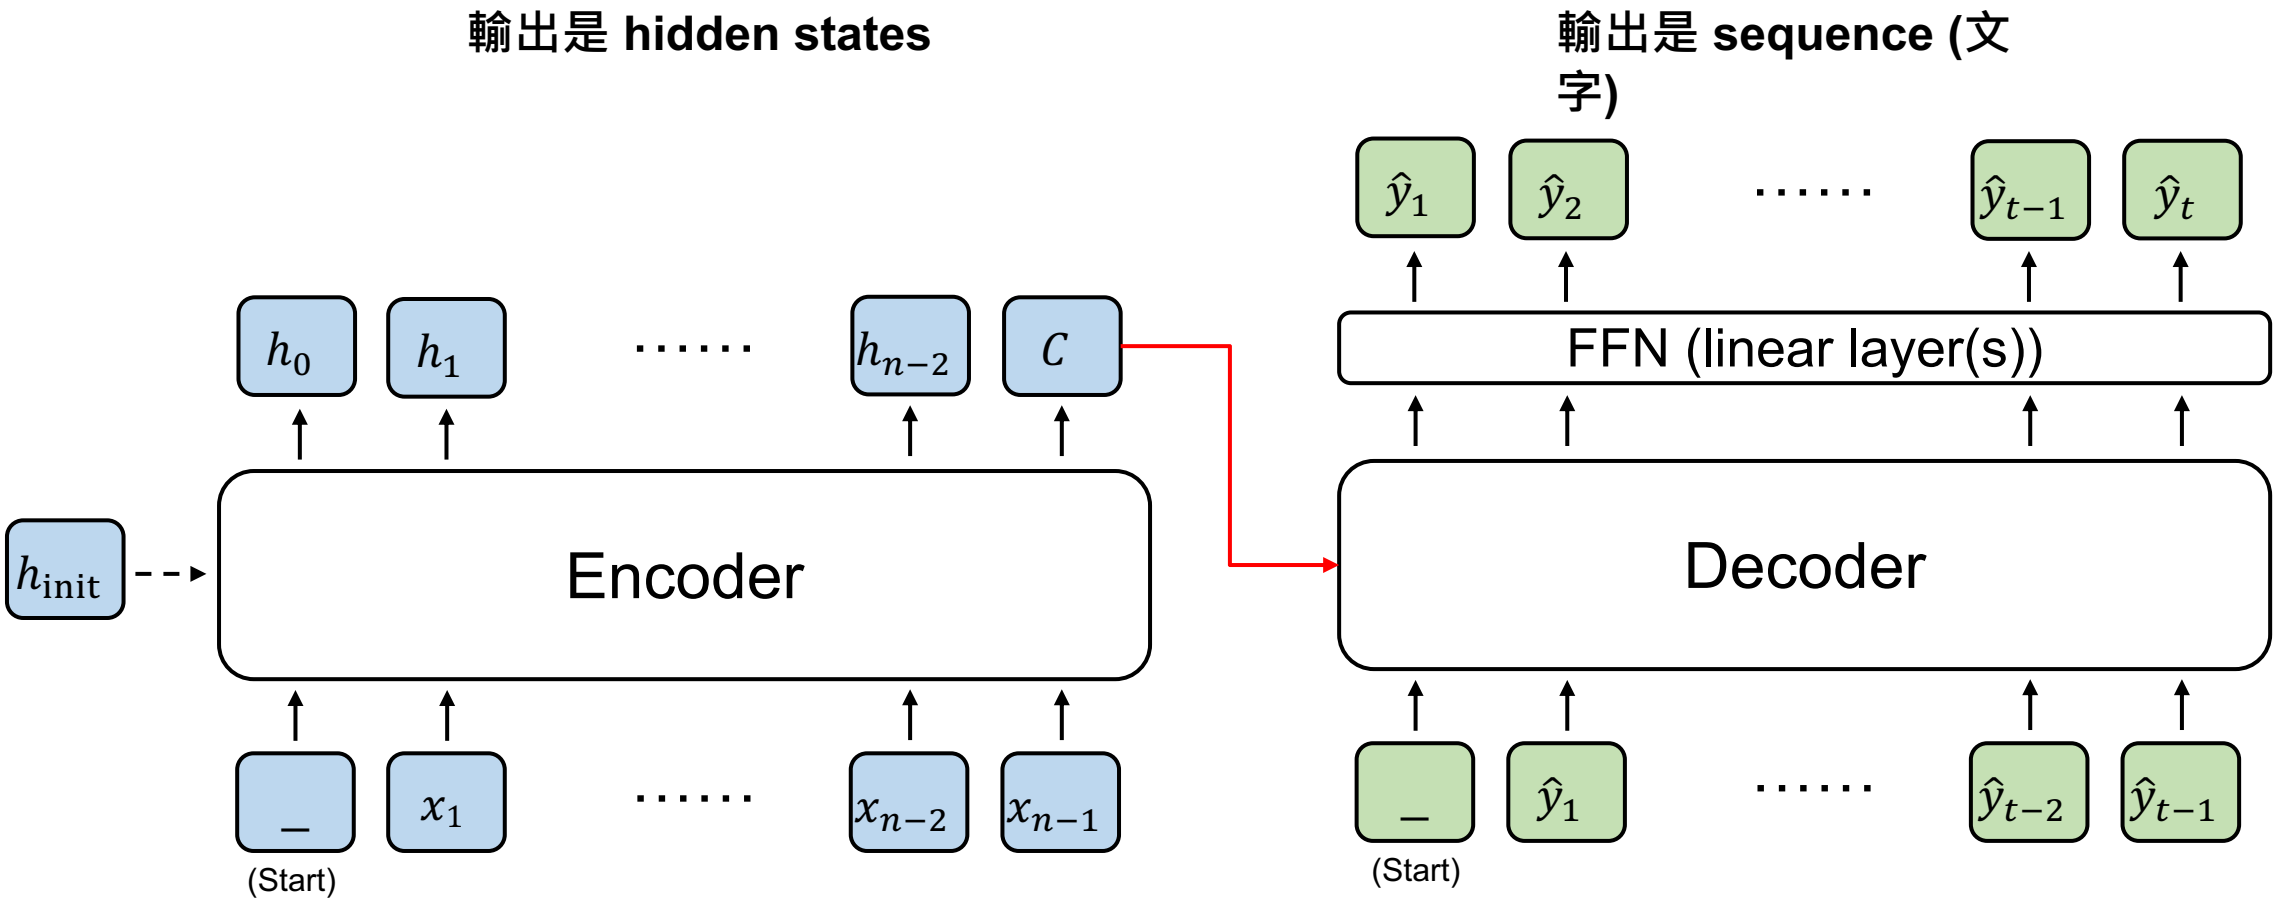

(Encoder 的輸出有很多種被Decoder運用的方式)

機器翻譯 (with Attention, t = 2)

$\overrightarrow{h_i} = \left[\overrightarrow{h_i}; \overleftarrow{h_i}\right]$

$$c_2 = \sum_{i=1}^{n} \alpha_{2i} h_t$$

i: encoder timestep
t: decoder timestep

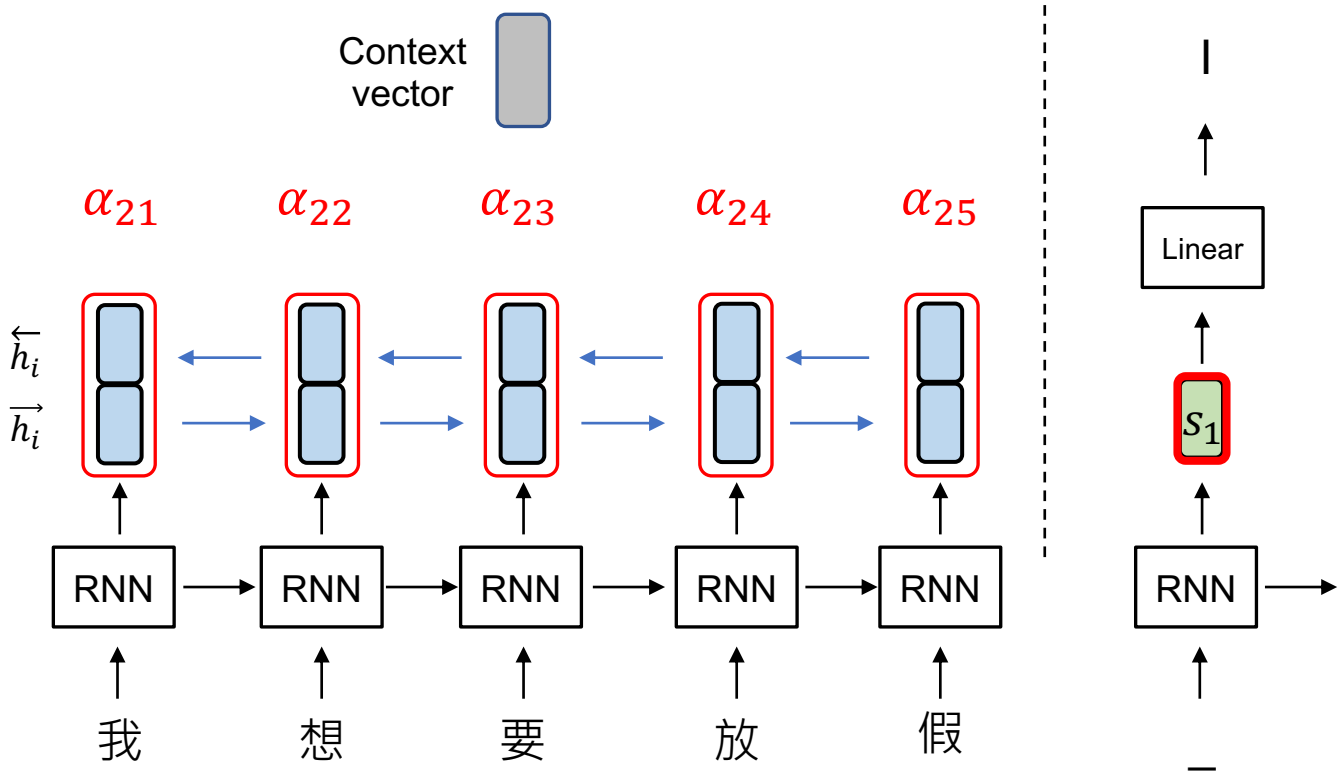

機器翻譯 (with Attention, t = 3)

$\overrightarrow{h_i} = \left[\overrightarrow{h_i}; \overleftarrow{h_i}\right]$

$$c_3 = \sum_{i=1}^{n} \alpha_{3i} h_t$$

i: encoder timestep
t: decoder timestep

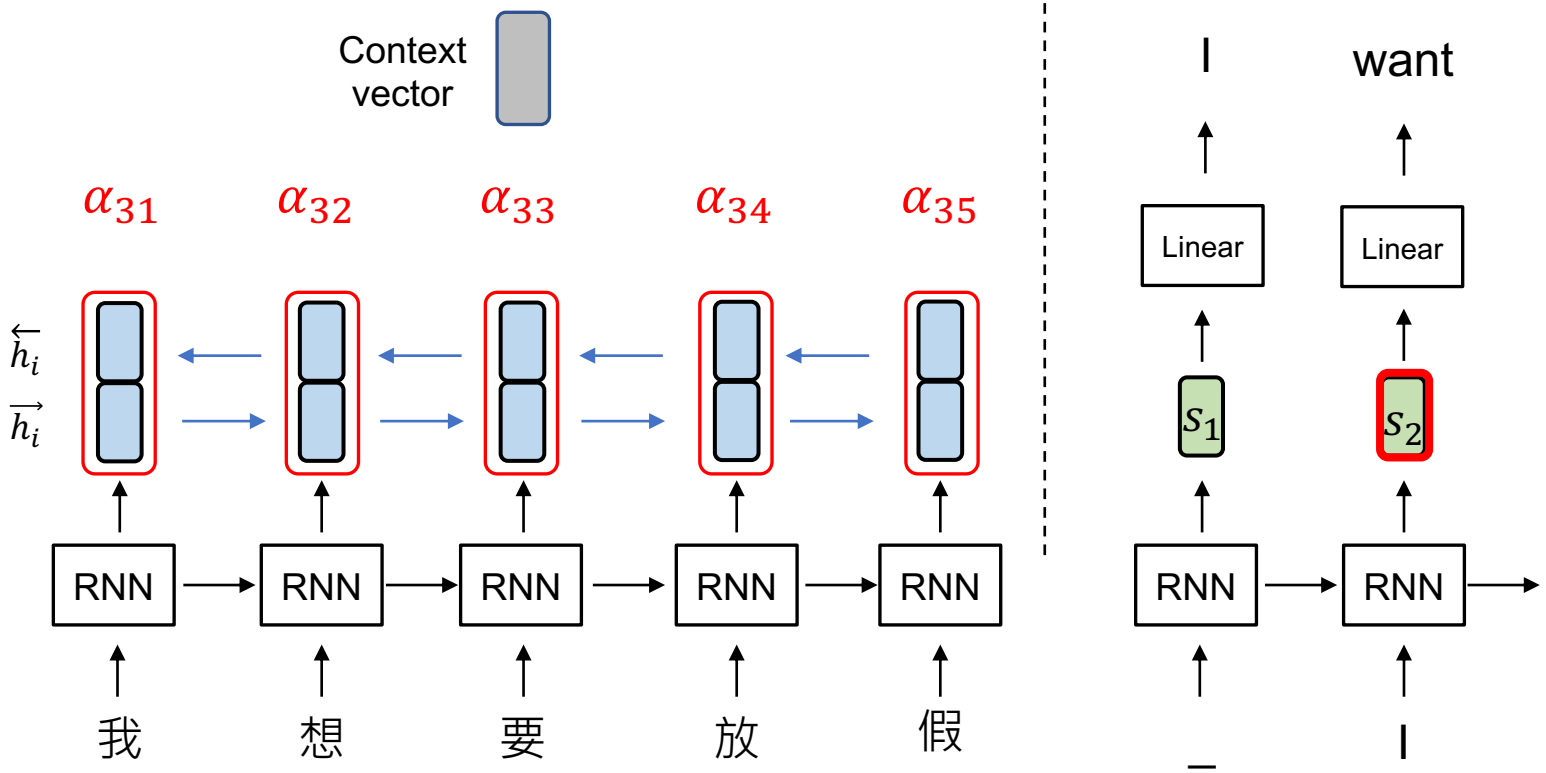

Additive Attention (算出 α_{ti})

$\overrightarrow{h_i} = \left[\overrightarrow{h_i}; \overleftarrow{h_i}\right]$

$$c_3 = \sum_{i=1}^{n} \alpha_{3i} h_t$$

i: encoder timestep
t: decoder timestep

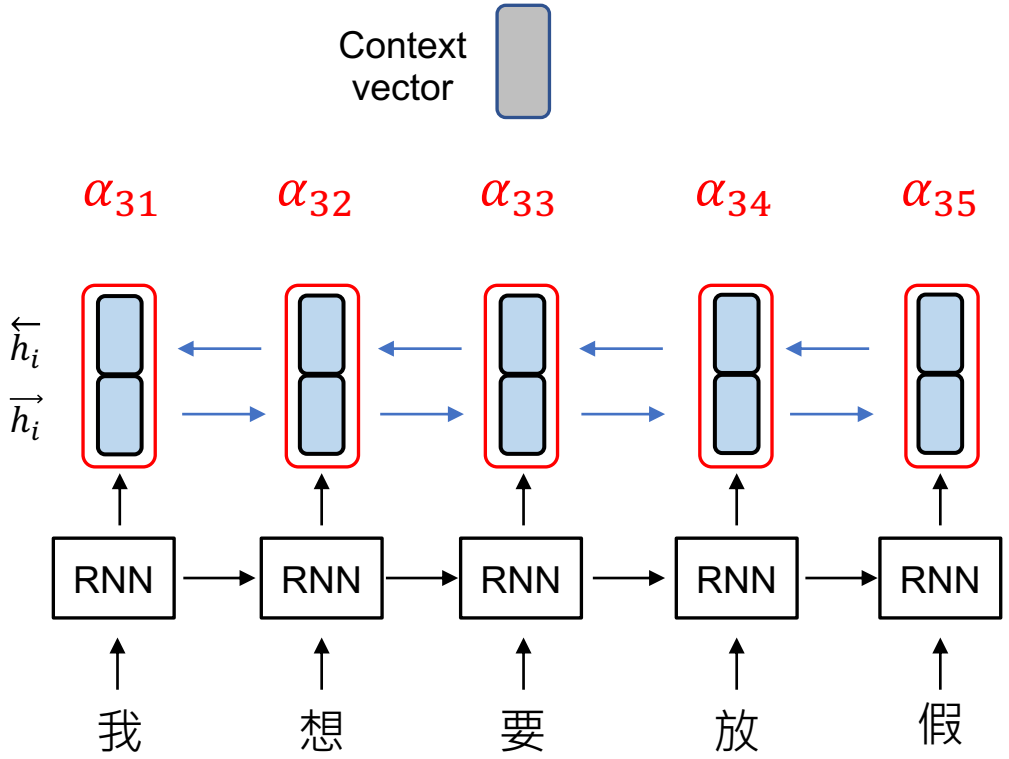

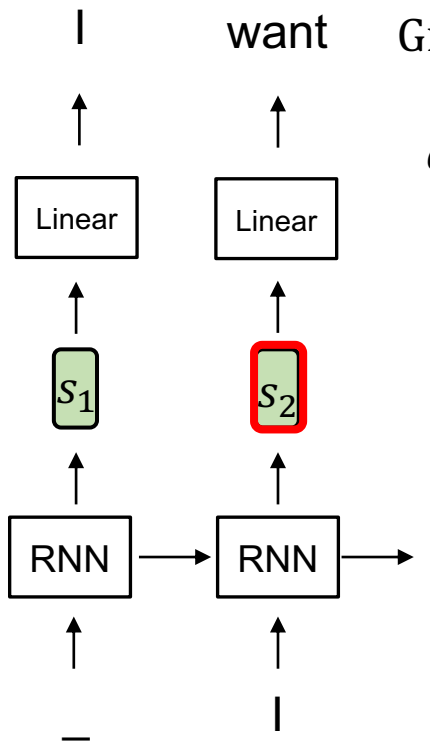

Given i ∈ {1, ..., n},

$$\alpha_{ti} = \text{softmax}(v \cdot \tanh(W s_{t-1} + U h_i))$$

Additive Attention

轉為1個數字

Transformer

RNN 與 CNN 作為序列模型的問題

問題：RNN 難以平行化

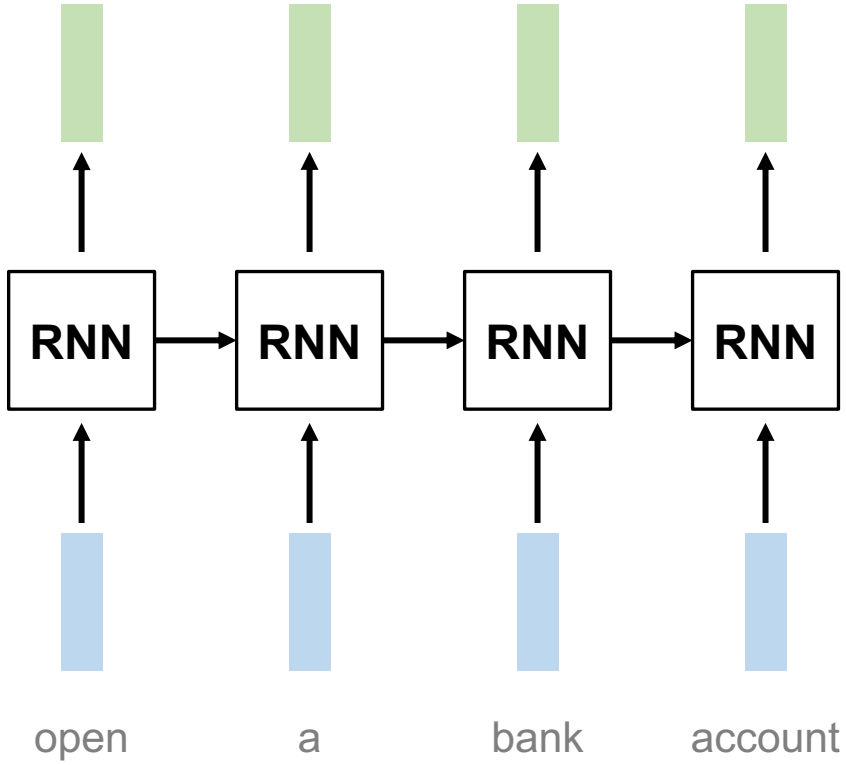

問題：CNN 著重局部資訊

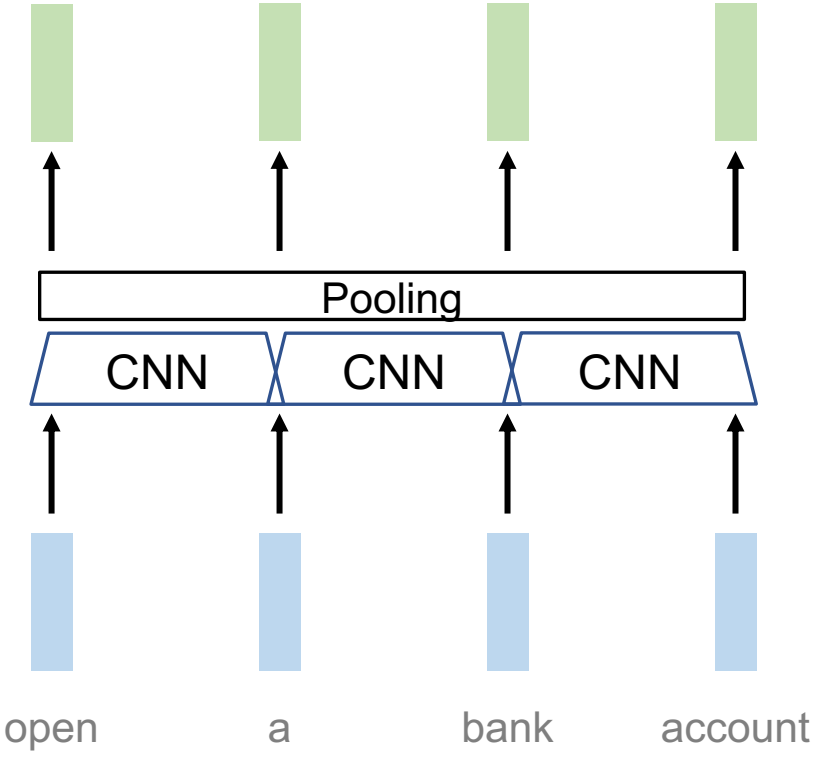

Overview of Transformer

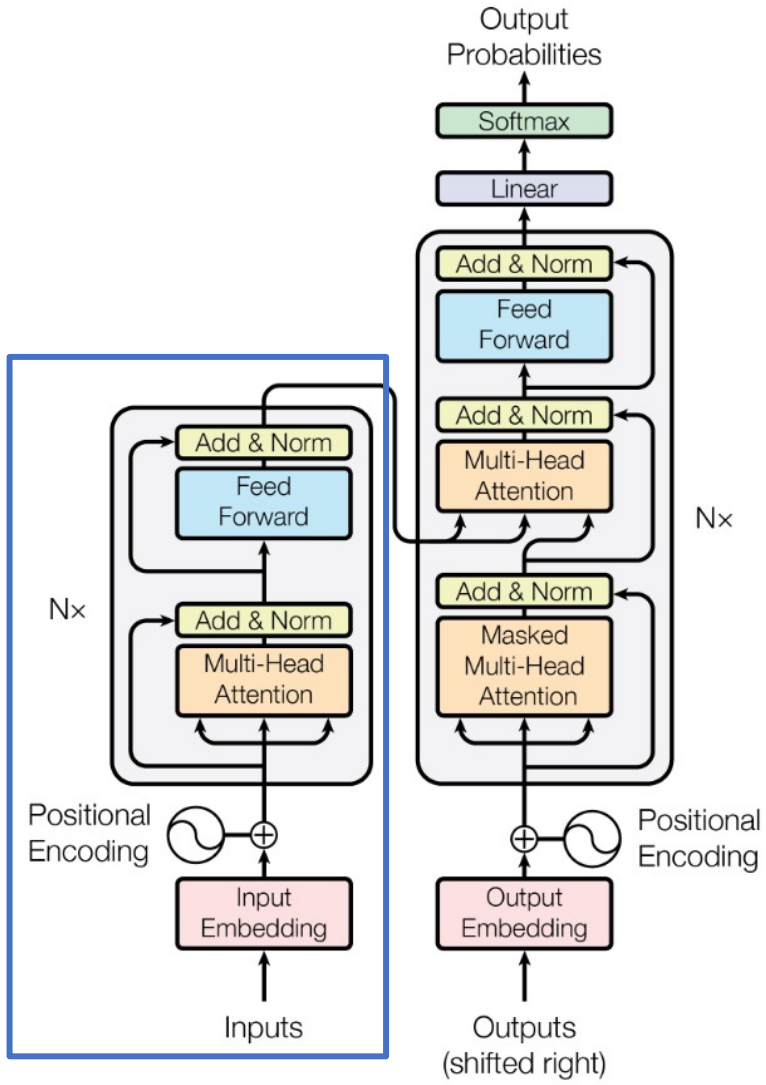

Encoder

Decoder

Vaswani, Ashish, et al. "Attention is all you need." NeurIPS (2017).

Outline of Transformer

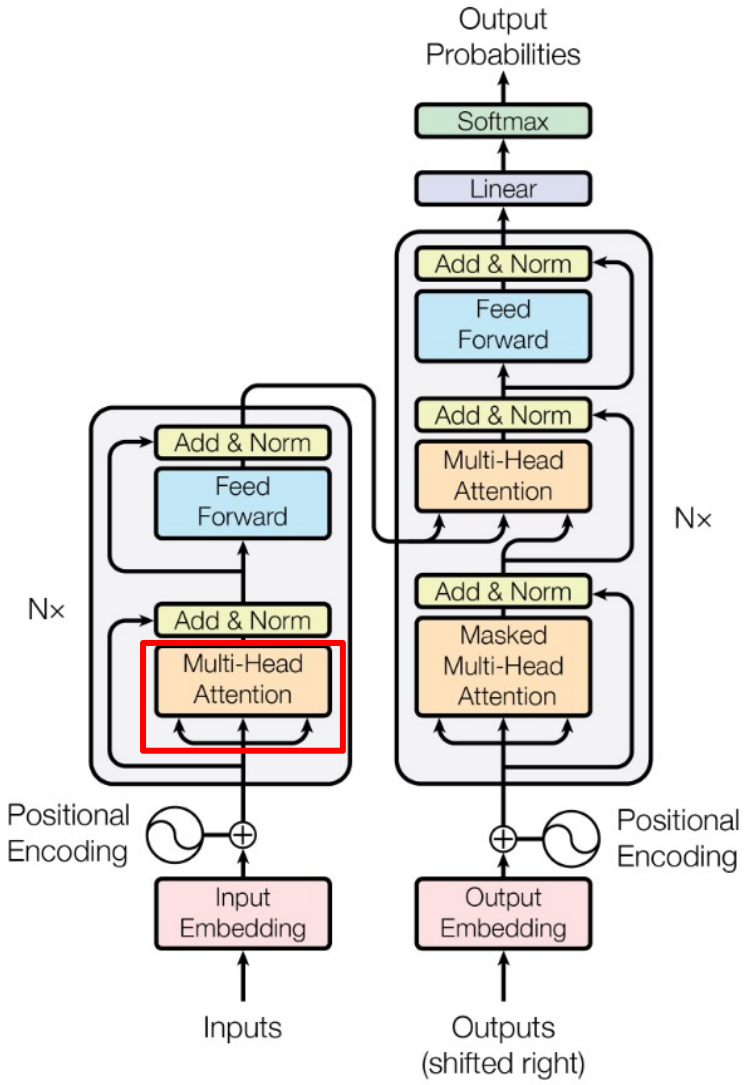

Vaswani, Ashish, et al. "Attention is all you need." NeurIPS (2017).

Transformer

Source: An example of the self-attention mechanism following long-distance... | Download Scientific Diagram (researchgate.net)

- Transformer 來自論文：Attention Is All You Need (Vaswani et al., 2017)
- Transformer 內部的主要機制為 Self-attention

Self-attention: 句子內的每個 token 自己跟自己做 attention

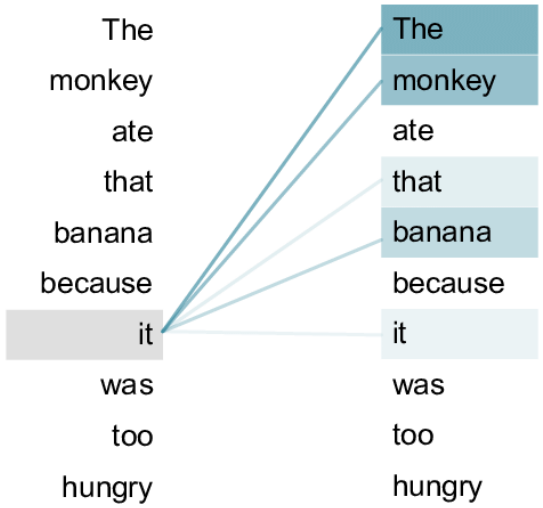

attention 進行的過程需要主動與被動

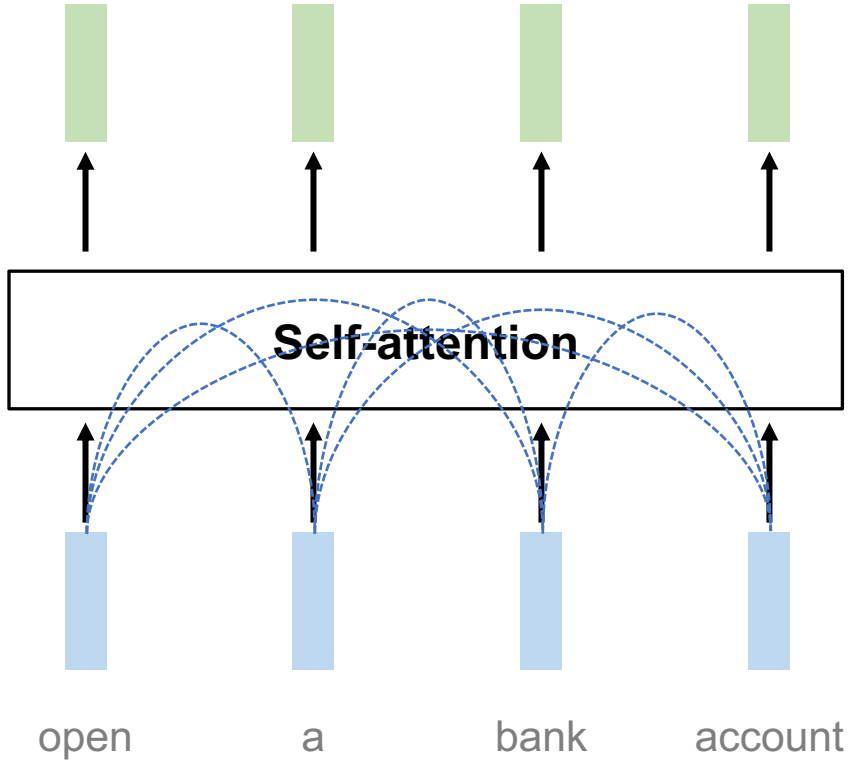

Query (Q), Key (K), and Value (V)

t = 1

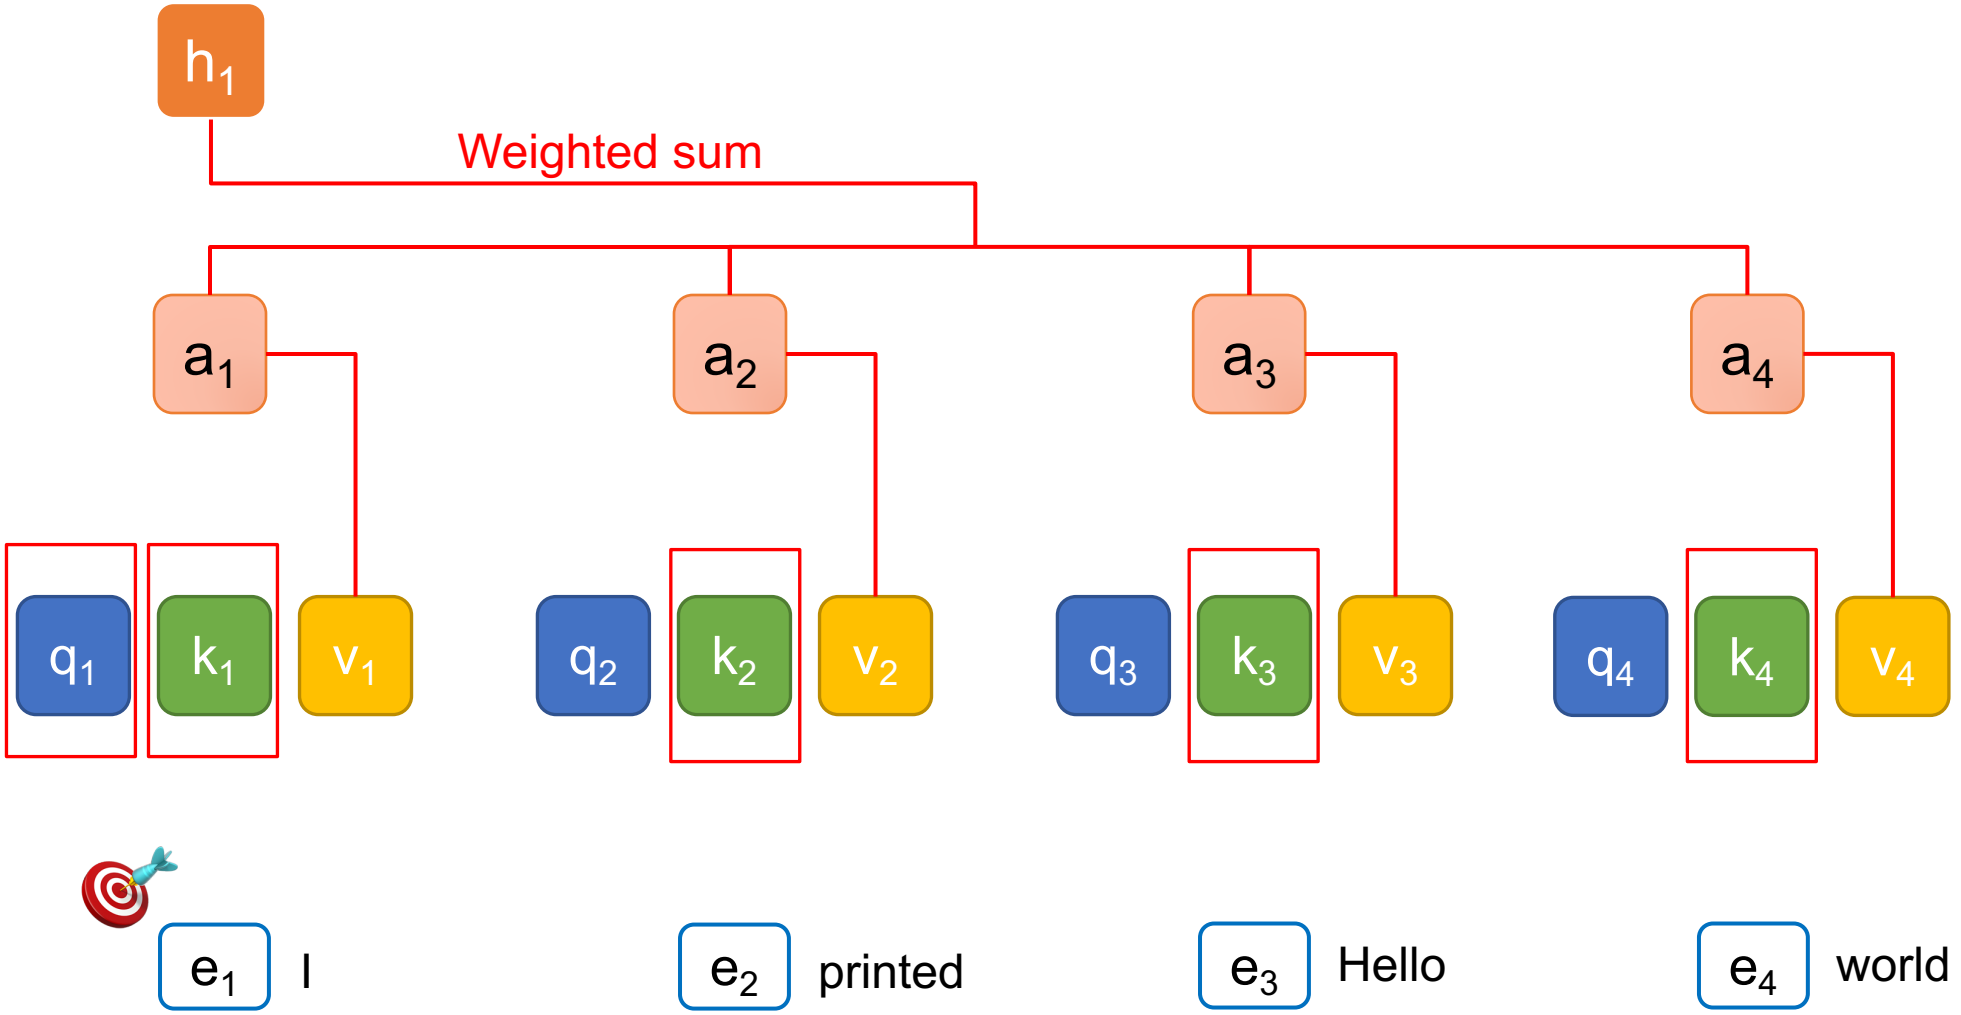

Query (Q), Key (K), and Value (V)

t = 2

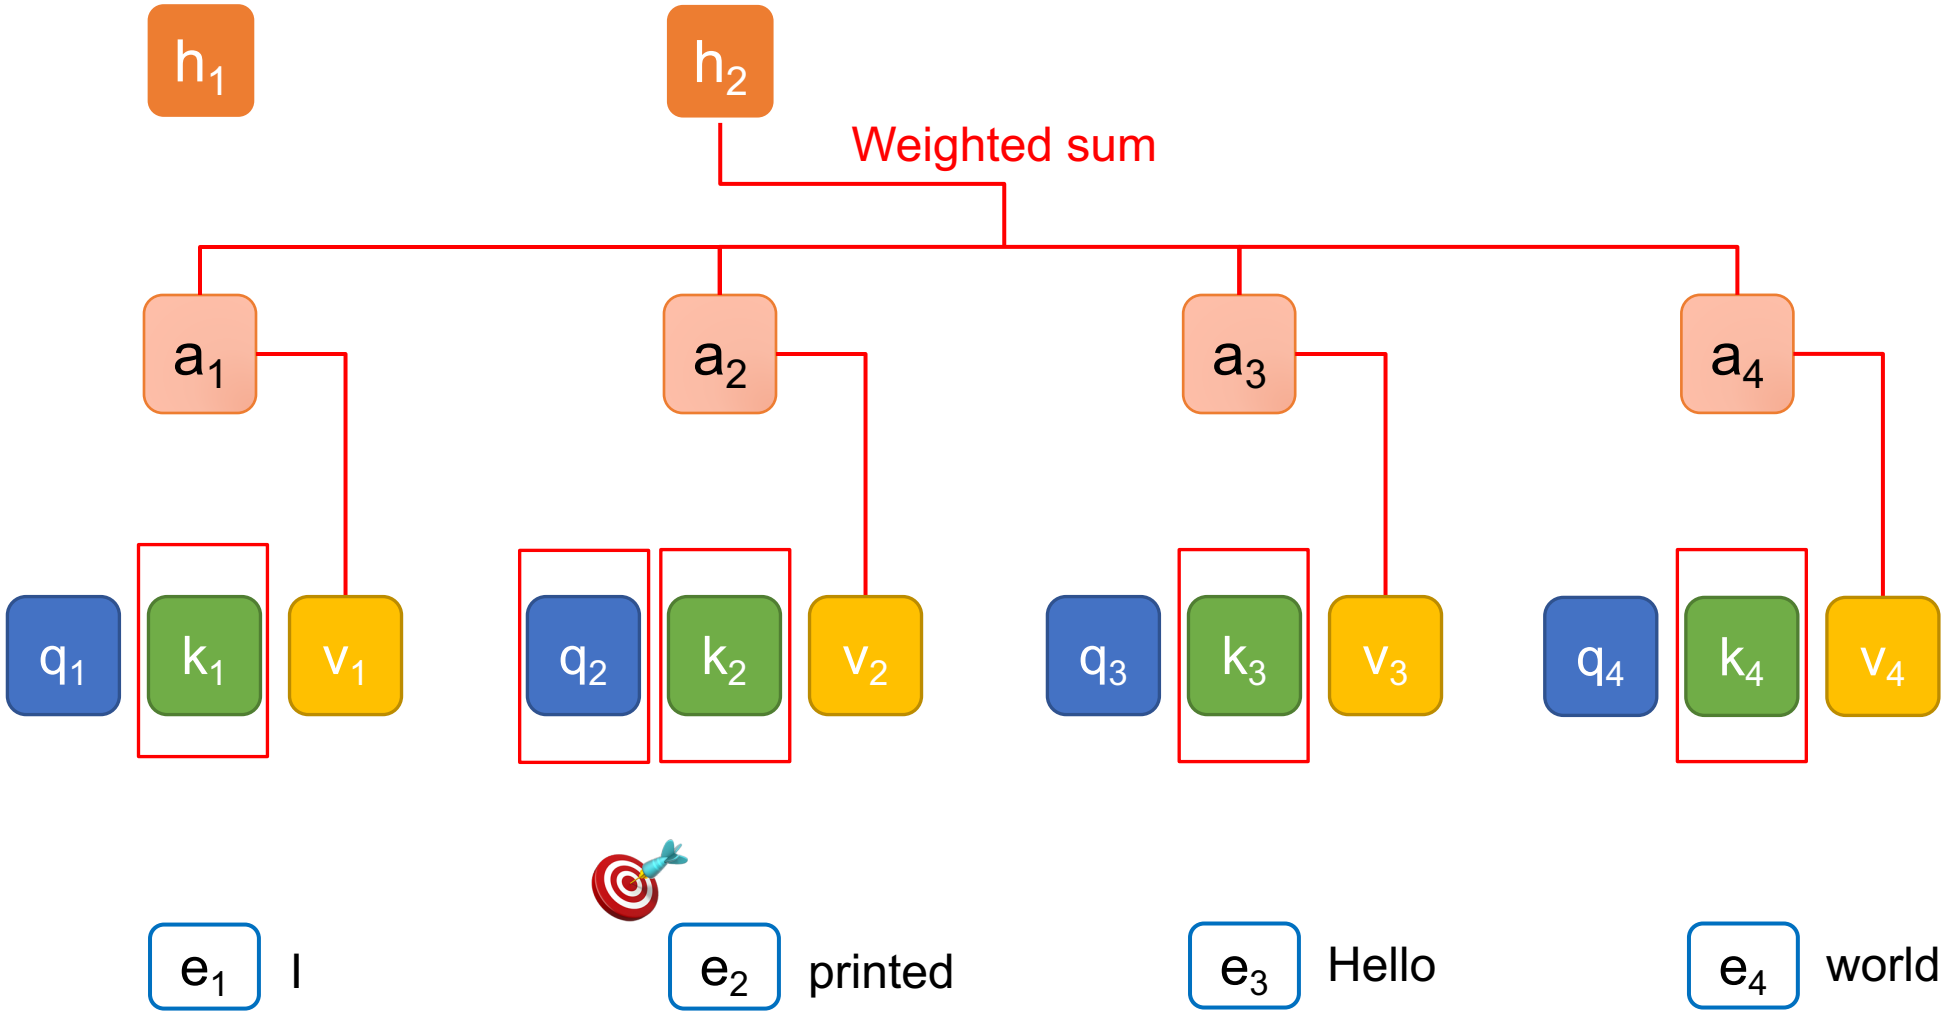

QKV 的比喻

Vaswani, Ashish, et al. "Attention is all you need." NeurIPS (2017).

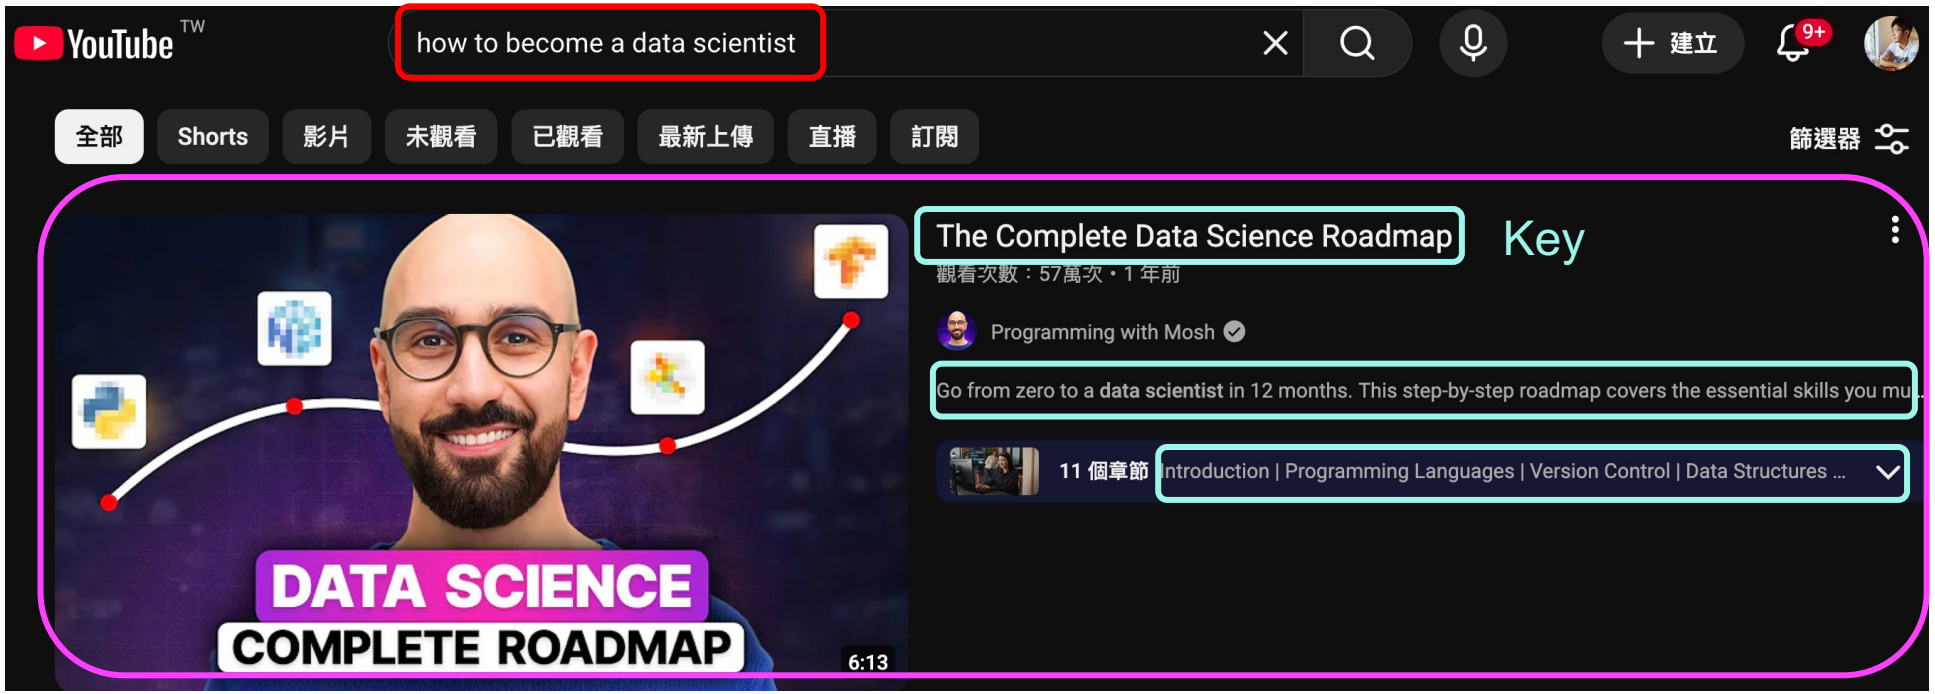

- Query: (主動) 查詢的關鍵字；Key: (被動) 查詢的對象
- Value: 值，關鍵字與對象的匹配程度

為什麼 Transformers 可以平行化？

d_k : queries 和 keys 的維度大小

通常設置 $d_q = d_k = d_v$

Scaled Dot-Product Attention

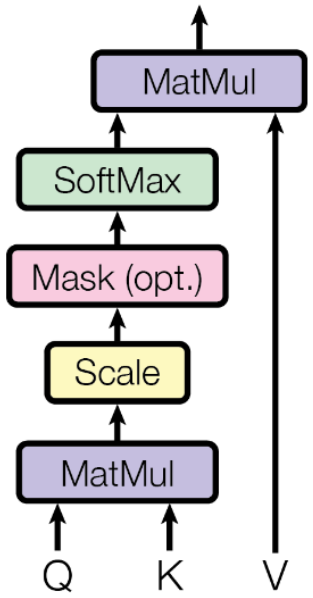

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

(自注意力機制可以利用矩陣乘積來進行平行化計算)

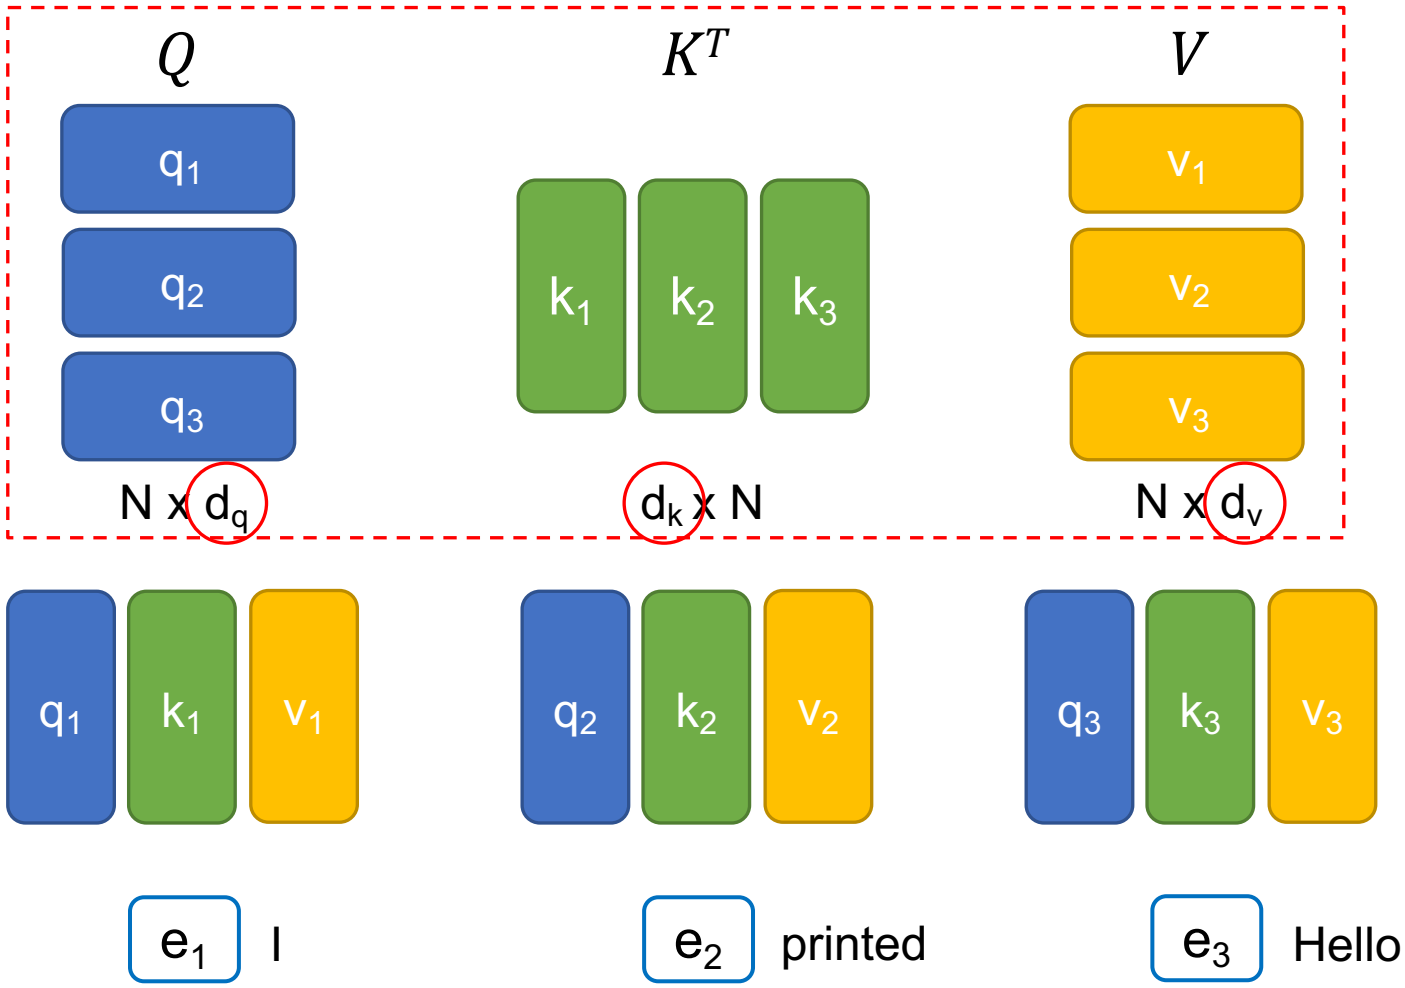

可以被訓練的參數

Transformers 的權重值 (1/2)

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

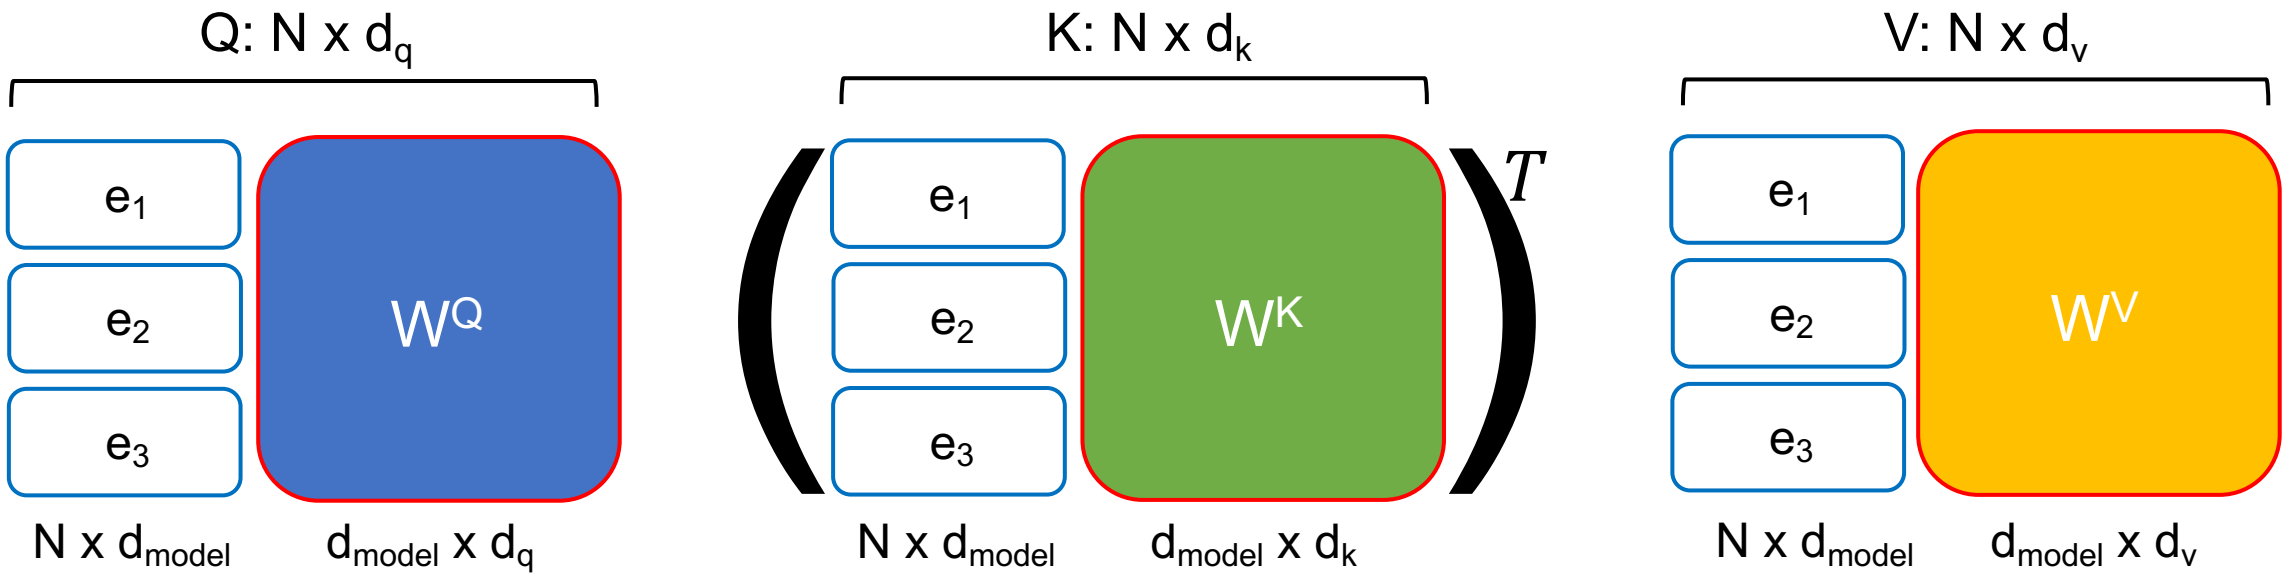

d_{model} : Embedding 的維度大小

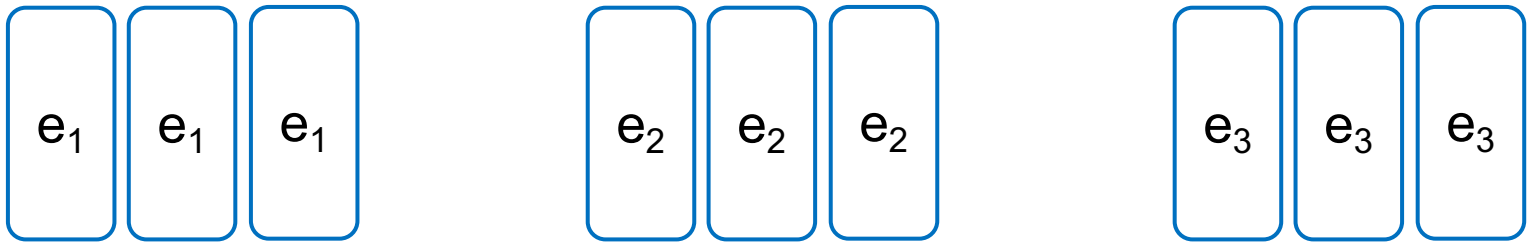

可以被訓練的參數

Transformers 的權重值 (2/2)

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

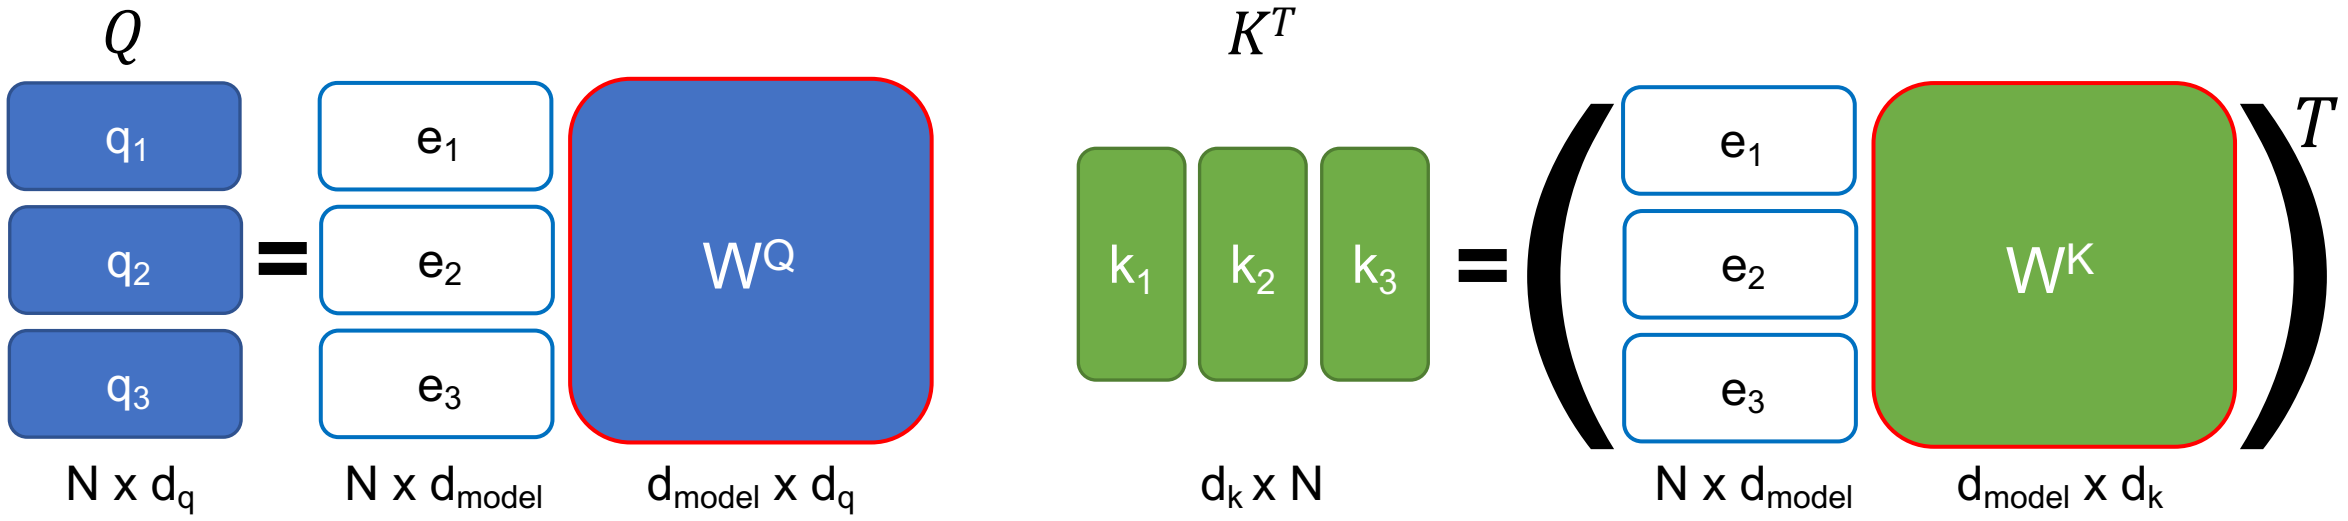

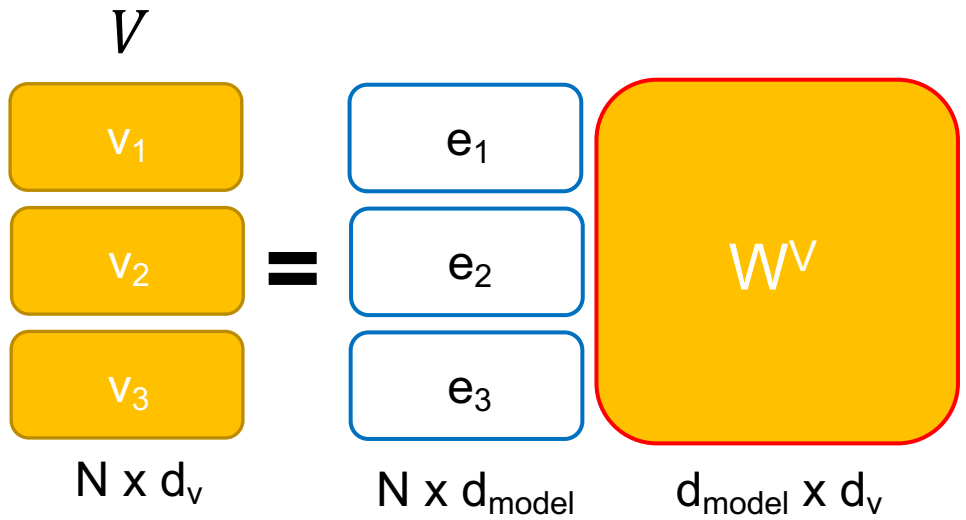

Self-attention 的公式與縮放項

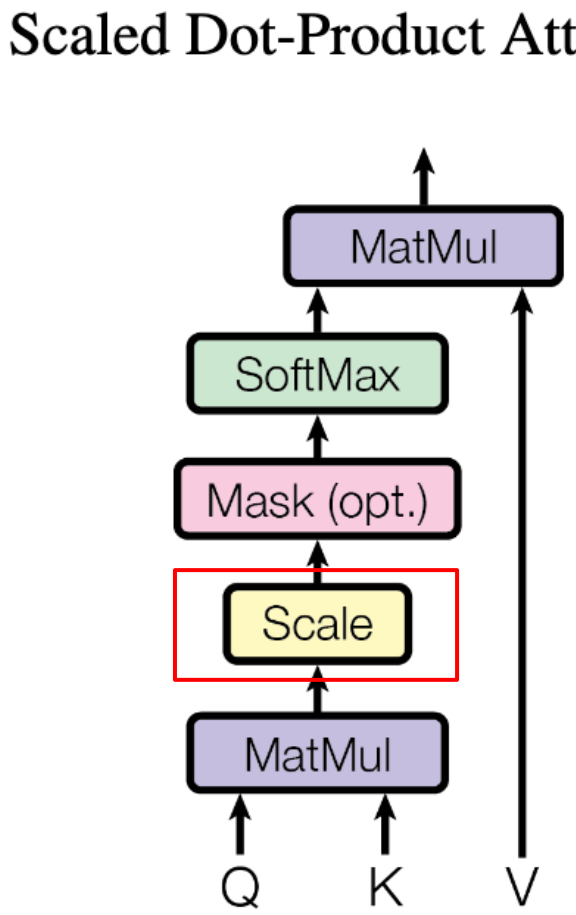

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

Query (Q), Key (K), and Value (V)

t = 1

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

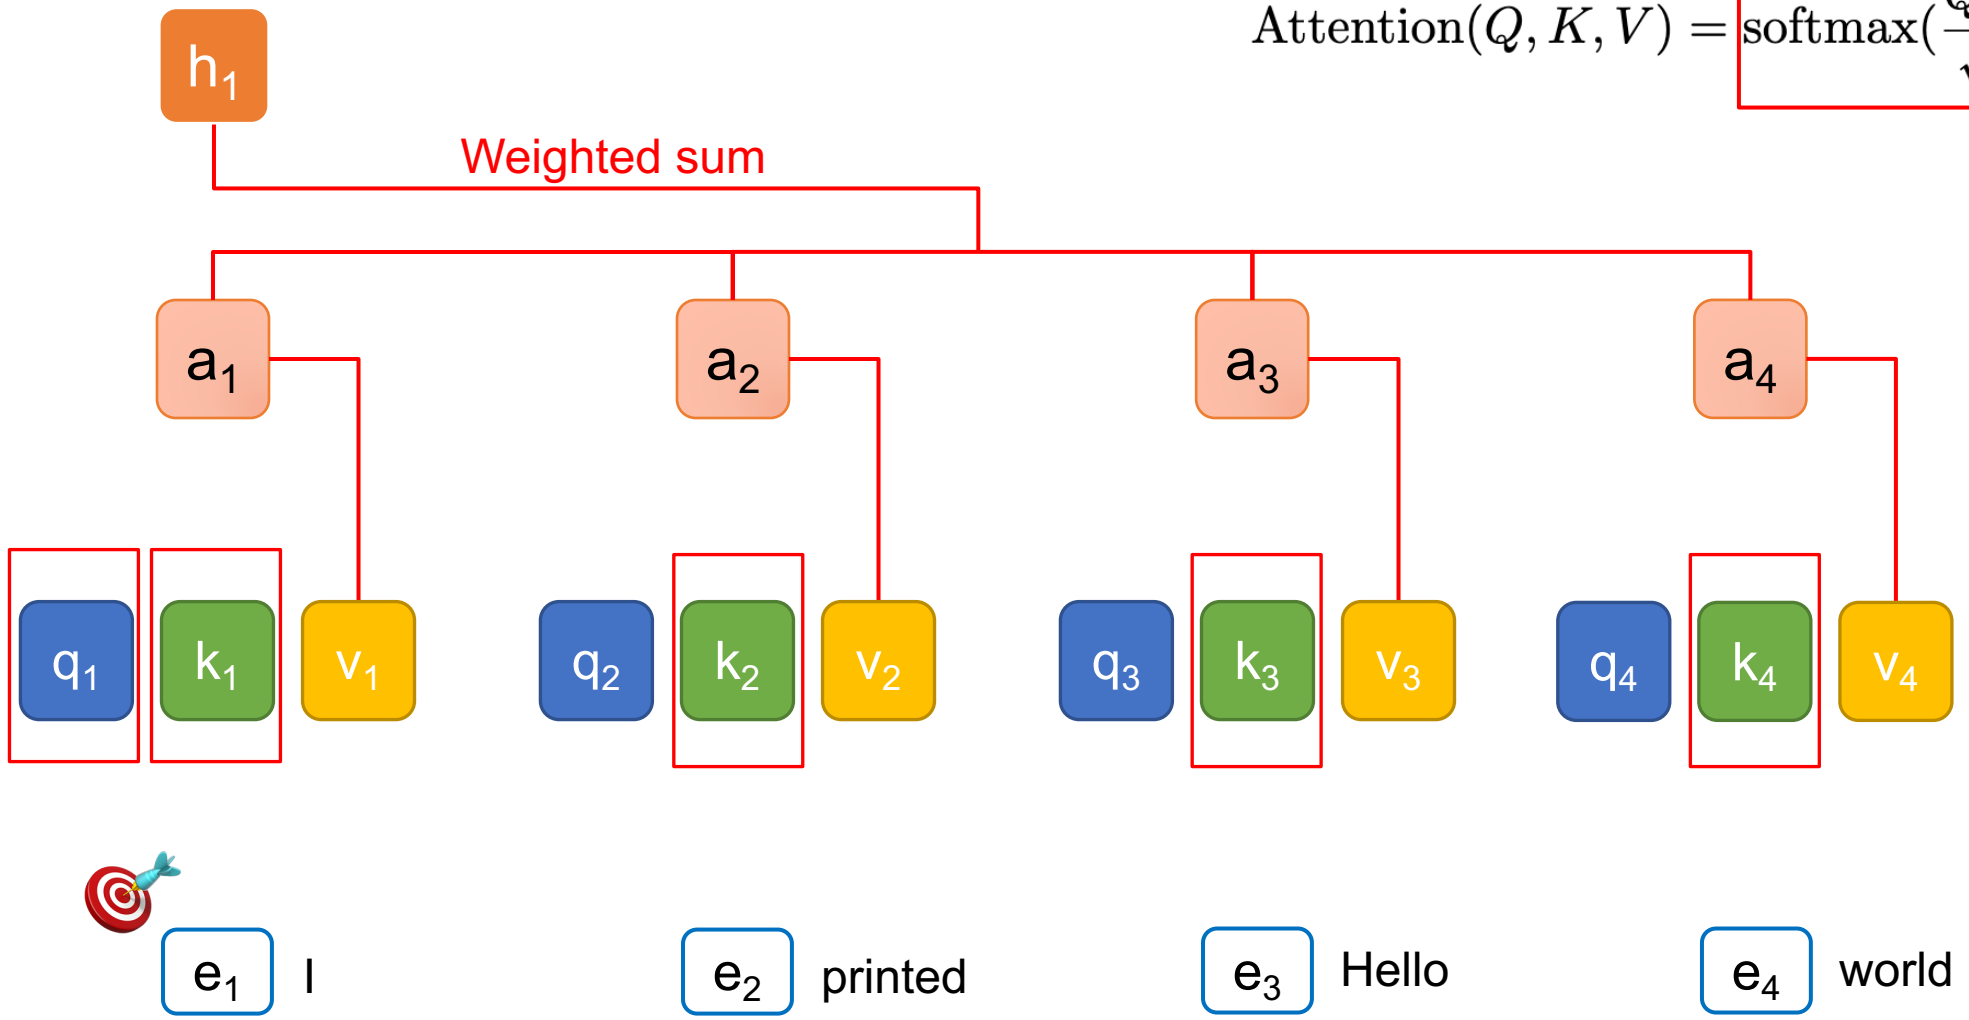

Query (Q), Key (K), and Value (V)

t = 2

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

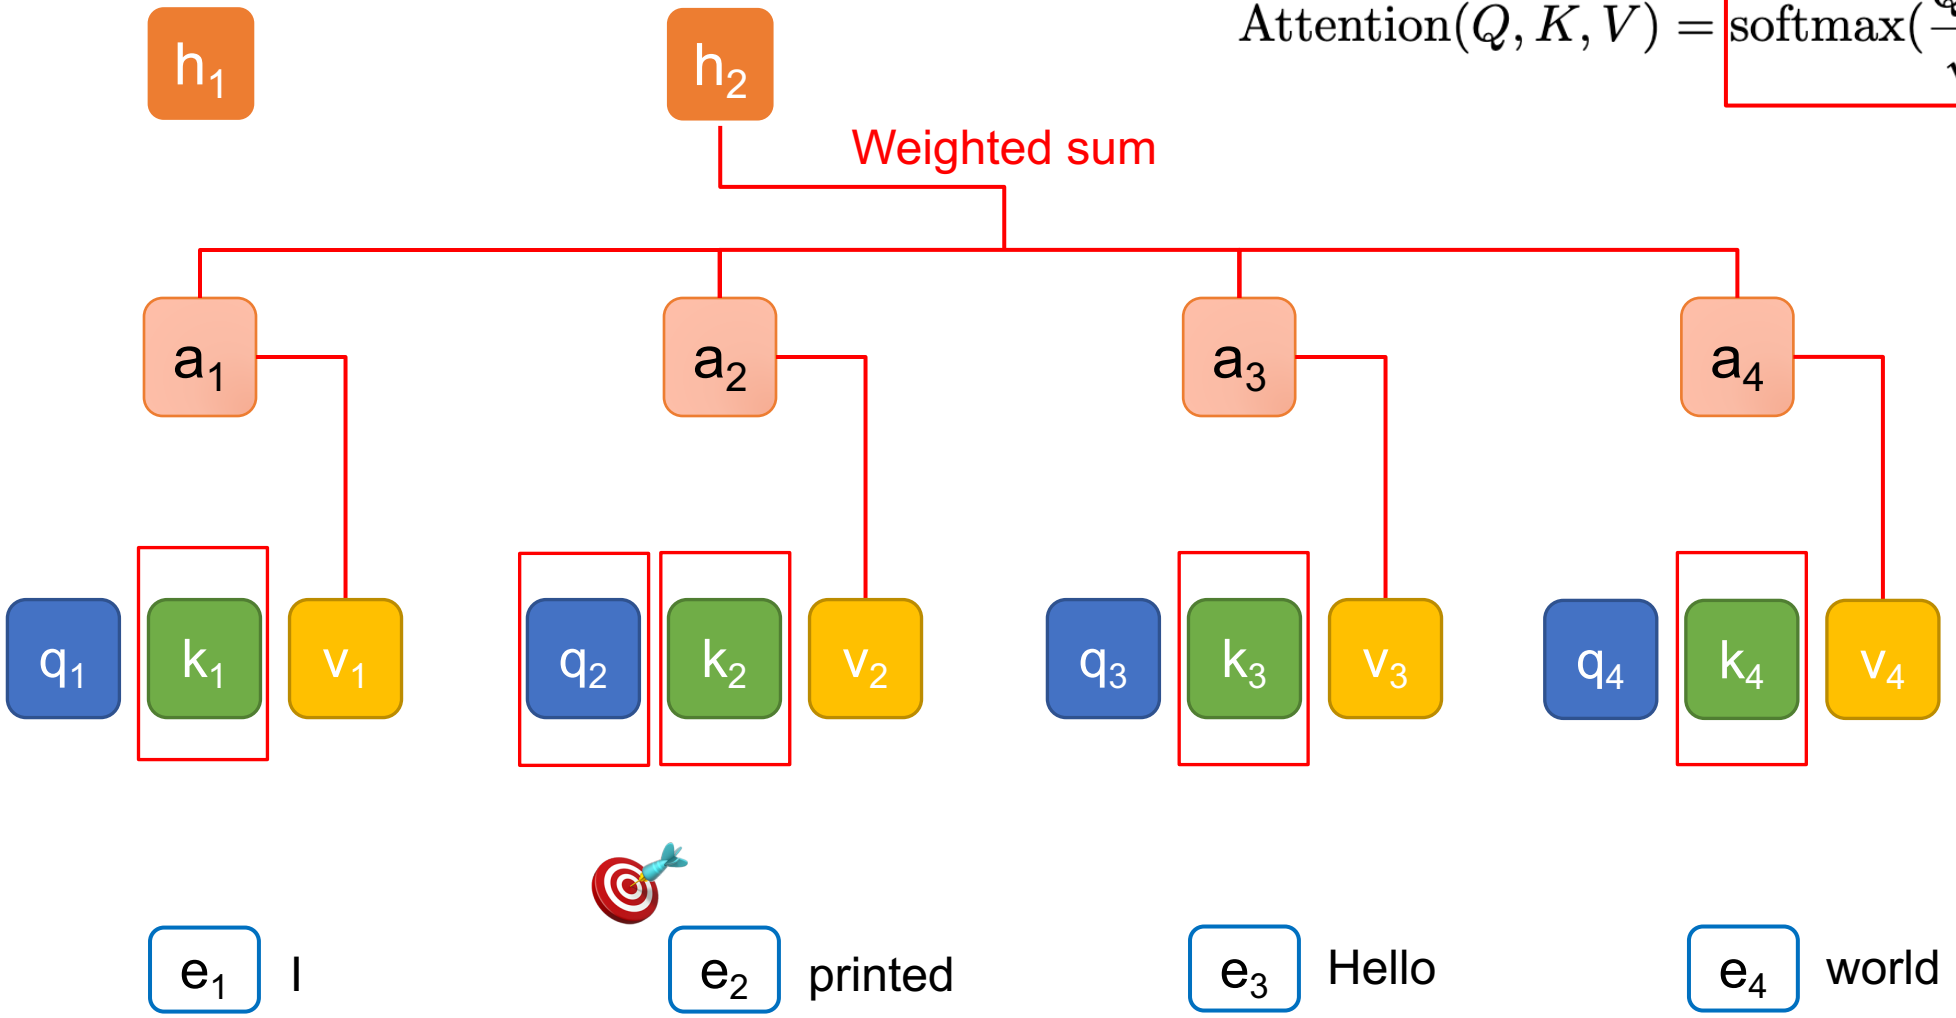

Self-Attention (Transformer) 過程

Vaswani et al. "Attention is all you need." NeurIPS 2017.
https://research.google/blog/transformer-a-novel-neural-network-architecture-for-language-understanding/

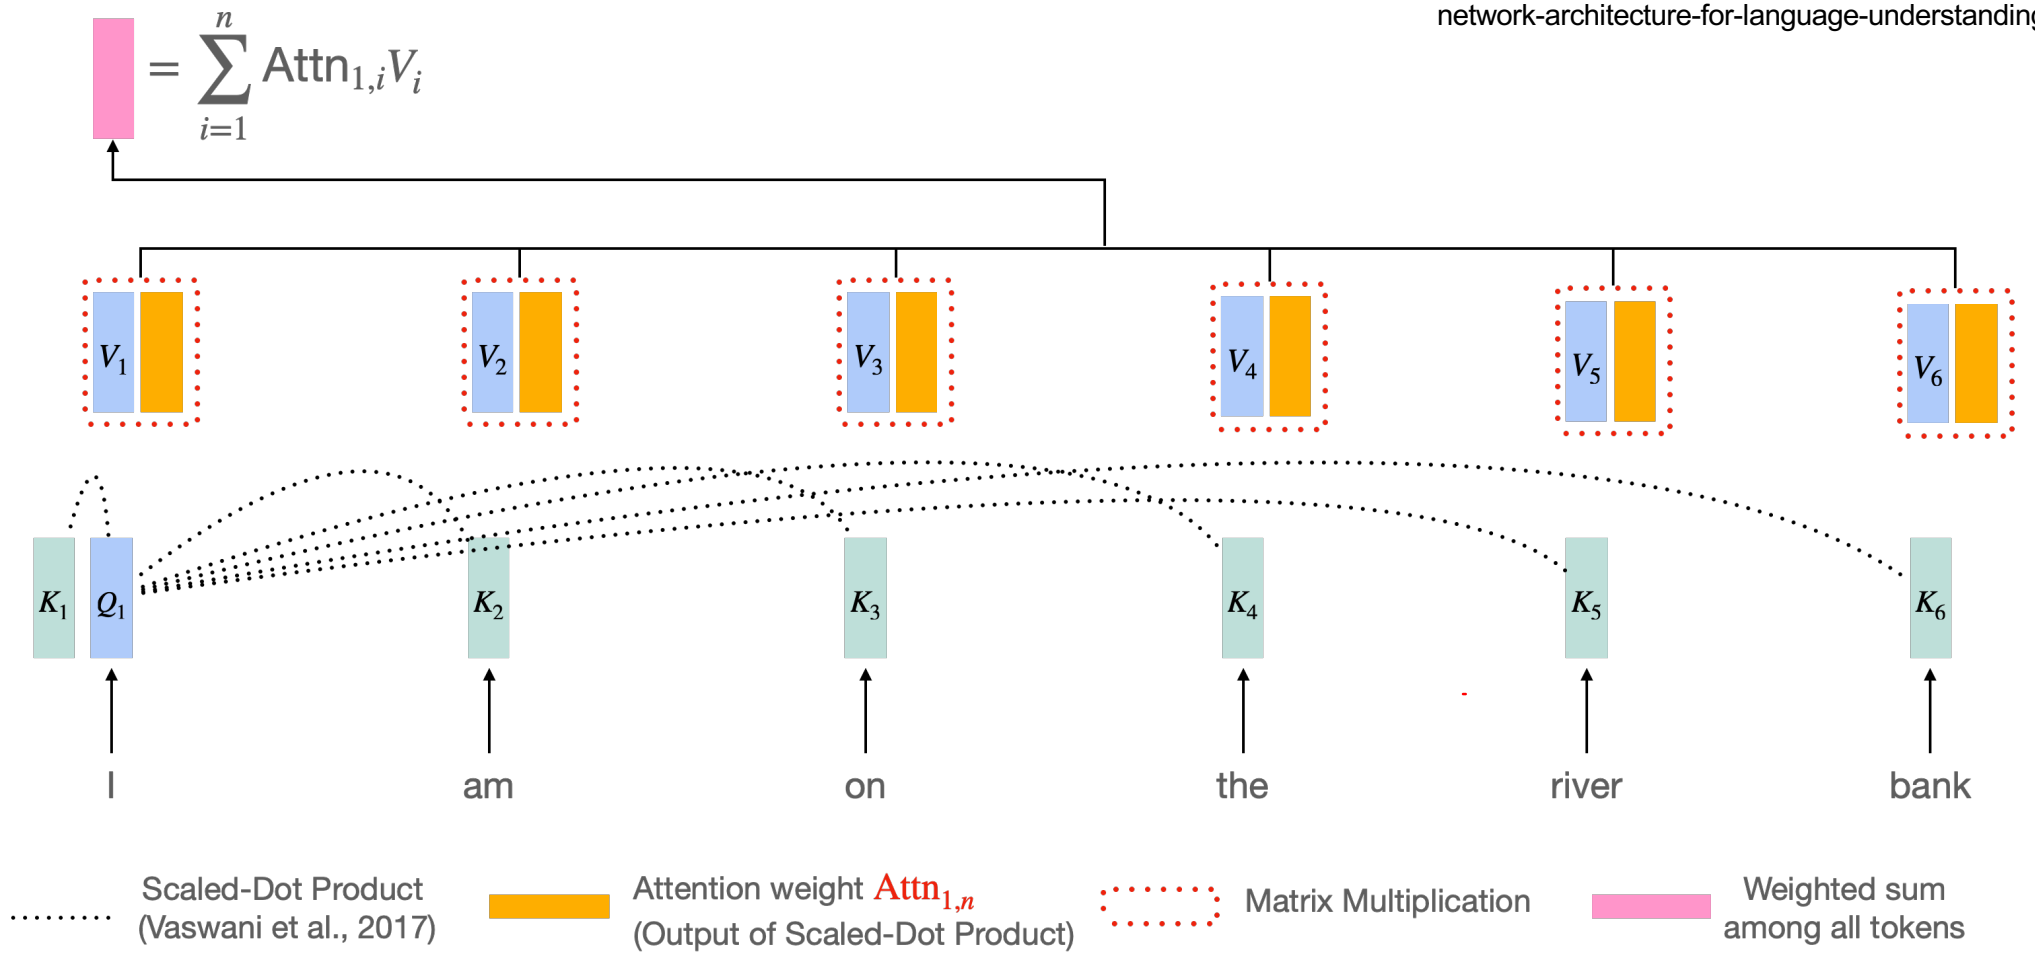

(For simplicity, I only draw the self-attention process for the first token.)

Self-Attention (Transformer) 過程

Vaswani et al. "Attention is all you need." NeurIPS 2017.
https://research.google/blog/transformer-a-novel-neural-network-architecture-for-language-understanding/

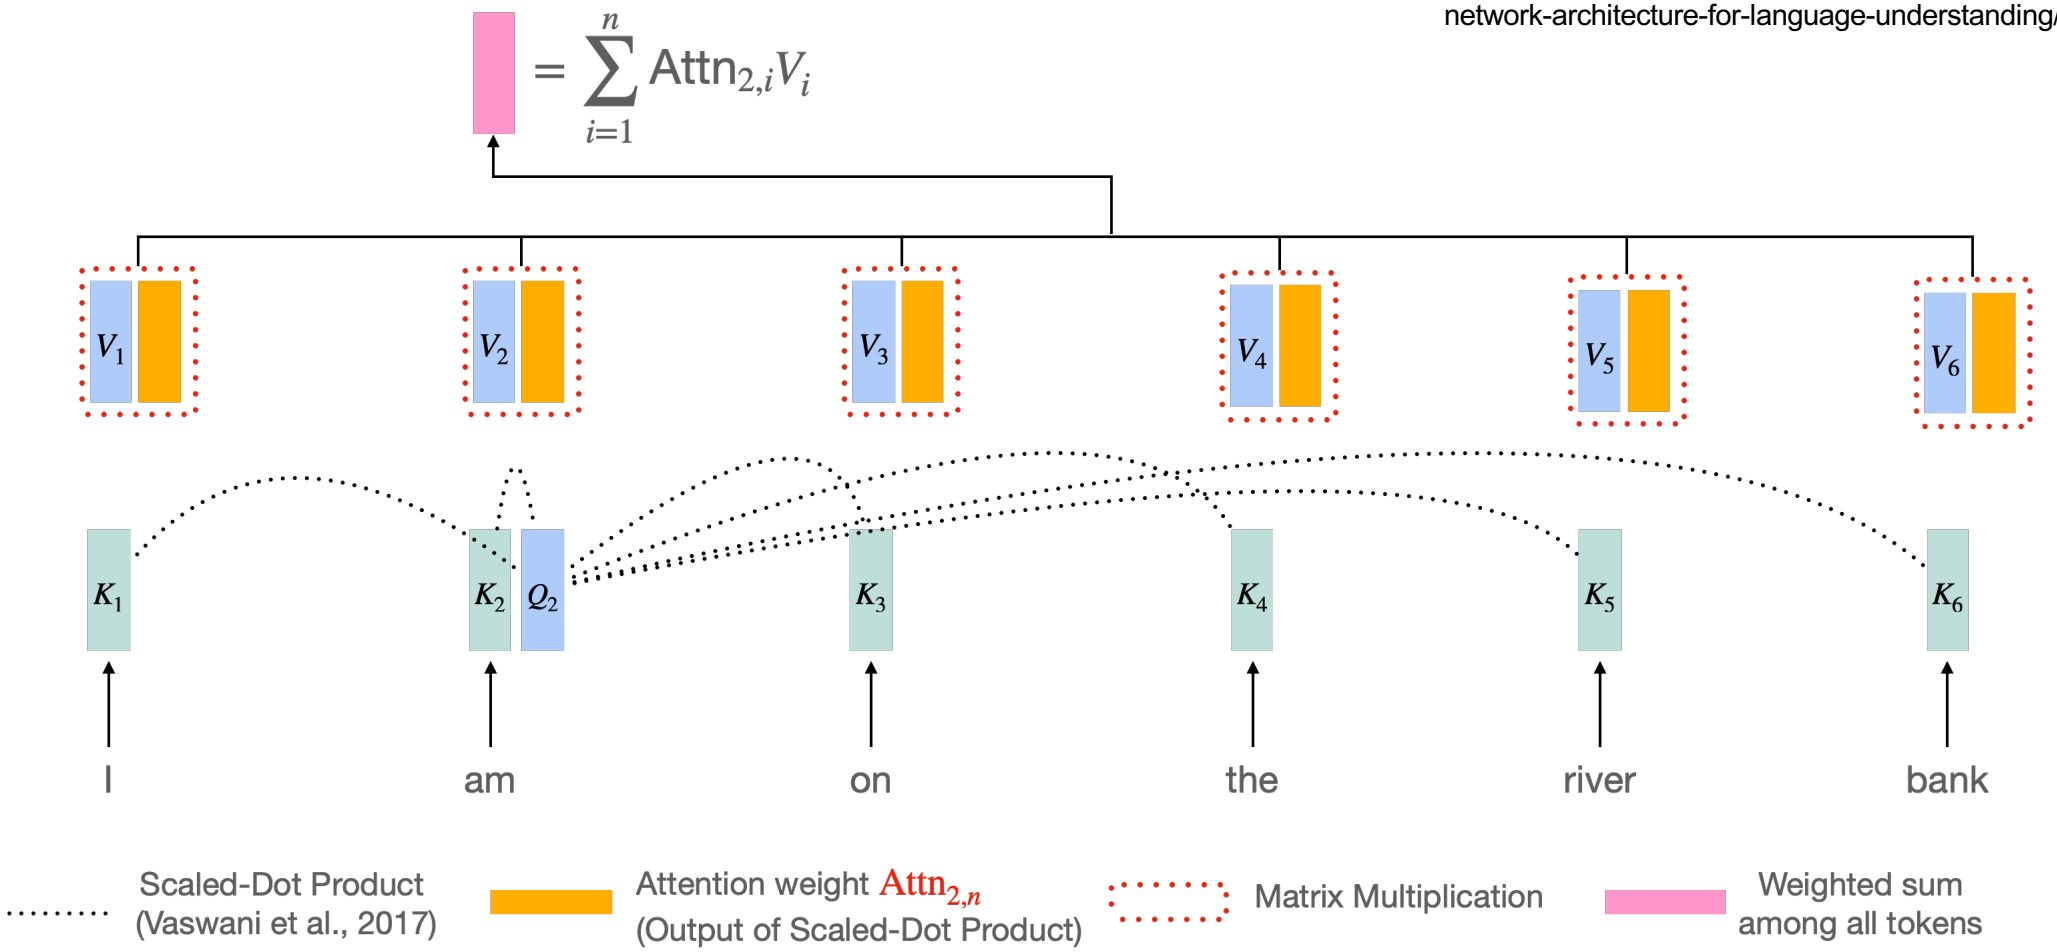

(For simplicity, I only draw the self-attention process for the second token.)

Self-attention 動畫

Figure source: https://research.google/blog/transformer-a-novel-neural-network-architecture-for-language-understanding/

Attention Is All You Need

Essential AI

Ashish Vaswani* Google Brain avaswani@google.com

Noam Shazeer* Google Brain noam@google.com

Anthropic

Niki Parmar* Google Research nikip@google.com

Inceptive

Jakob Uszkoreit* Google Research usz@google.com

Cohere

Sakana AI

Llion Jones* Google Research llion@google.com

Aidan N. Gomez* † University of Toronto aidan@cs.toronto.edu

Łukasz Kaiser* Google Brain lukaszkaiser@google.com

OpenAI

NEAR

Illia Polosukhin* ‡ illia.polosukhin@gmail.com

https://arxiv.org/abs/1706.03762v6

Multi-Head Attention (MHA)

- 到目前為止，我們已經介紹了 self-attention
  - We query the search engine
  - We search engine tries to map our query to the keys
  - The outputs are the attended values (e.g., the best matched videos).
- 試想：如果我們想要模型注意到多種語言行爲呢？

語言存在多種行為 / 模式

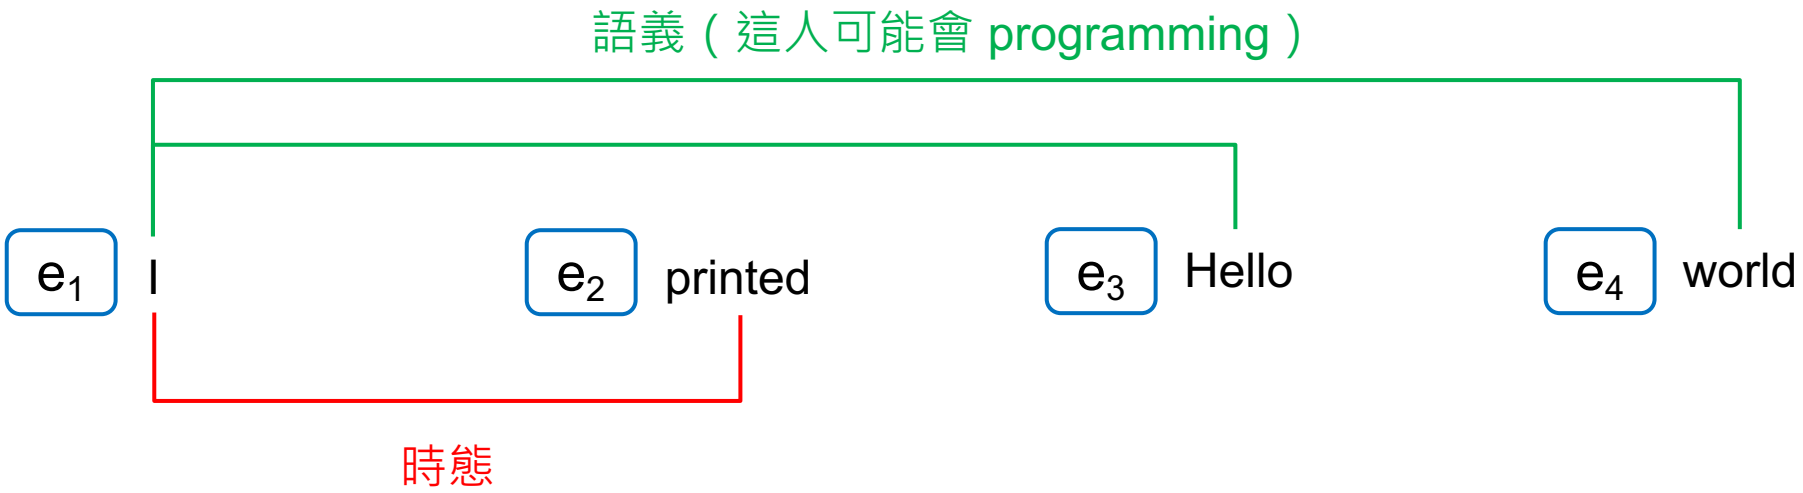

h_1 <- Softmax 也是一種 max，只能取得一種最有可能的語言模式

有沒有可能讓模型同時學到多種語言模式？
-> Multi-head Attention

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

Multi-Head Attention (MHA)

上標: head
下標: token index

Head = 1 (without MHA)

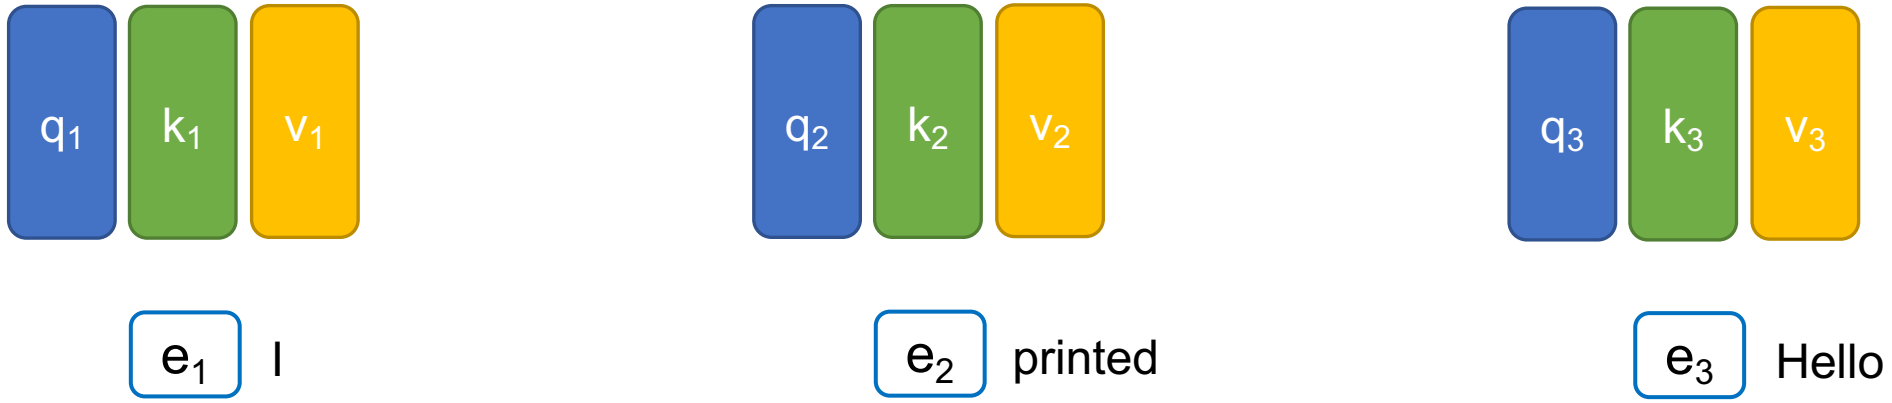

Head = 2 (with MHA)

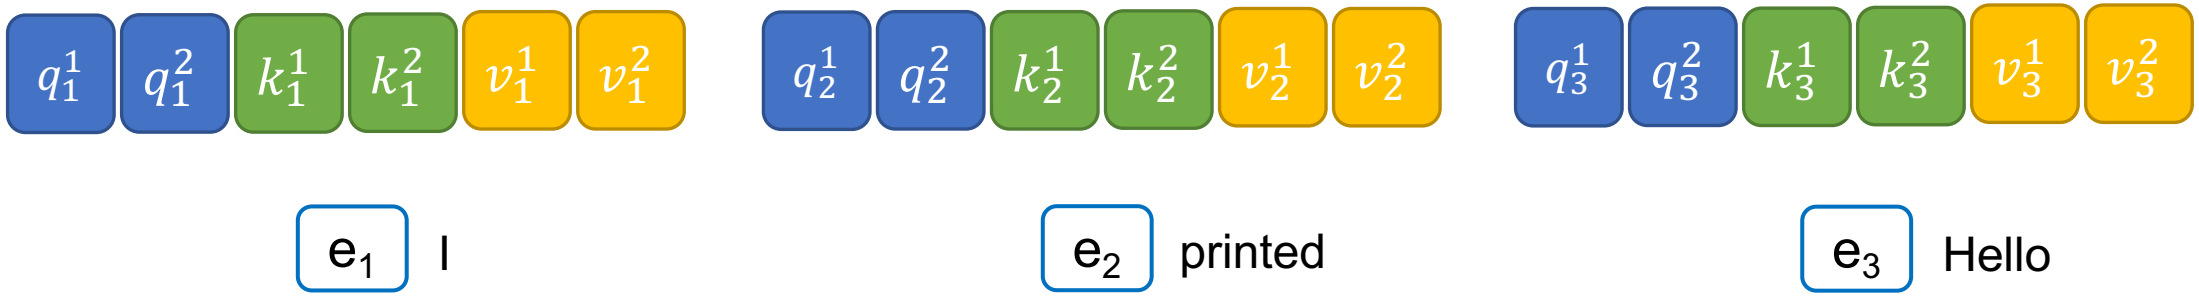

Query (Q), Key (K), and Value (V) with MHA

Head = 2, 使用第一個 head

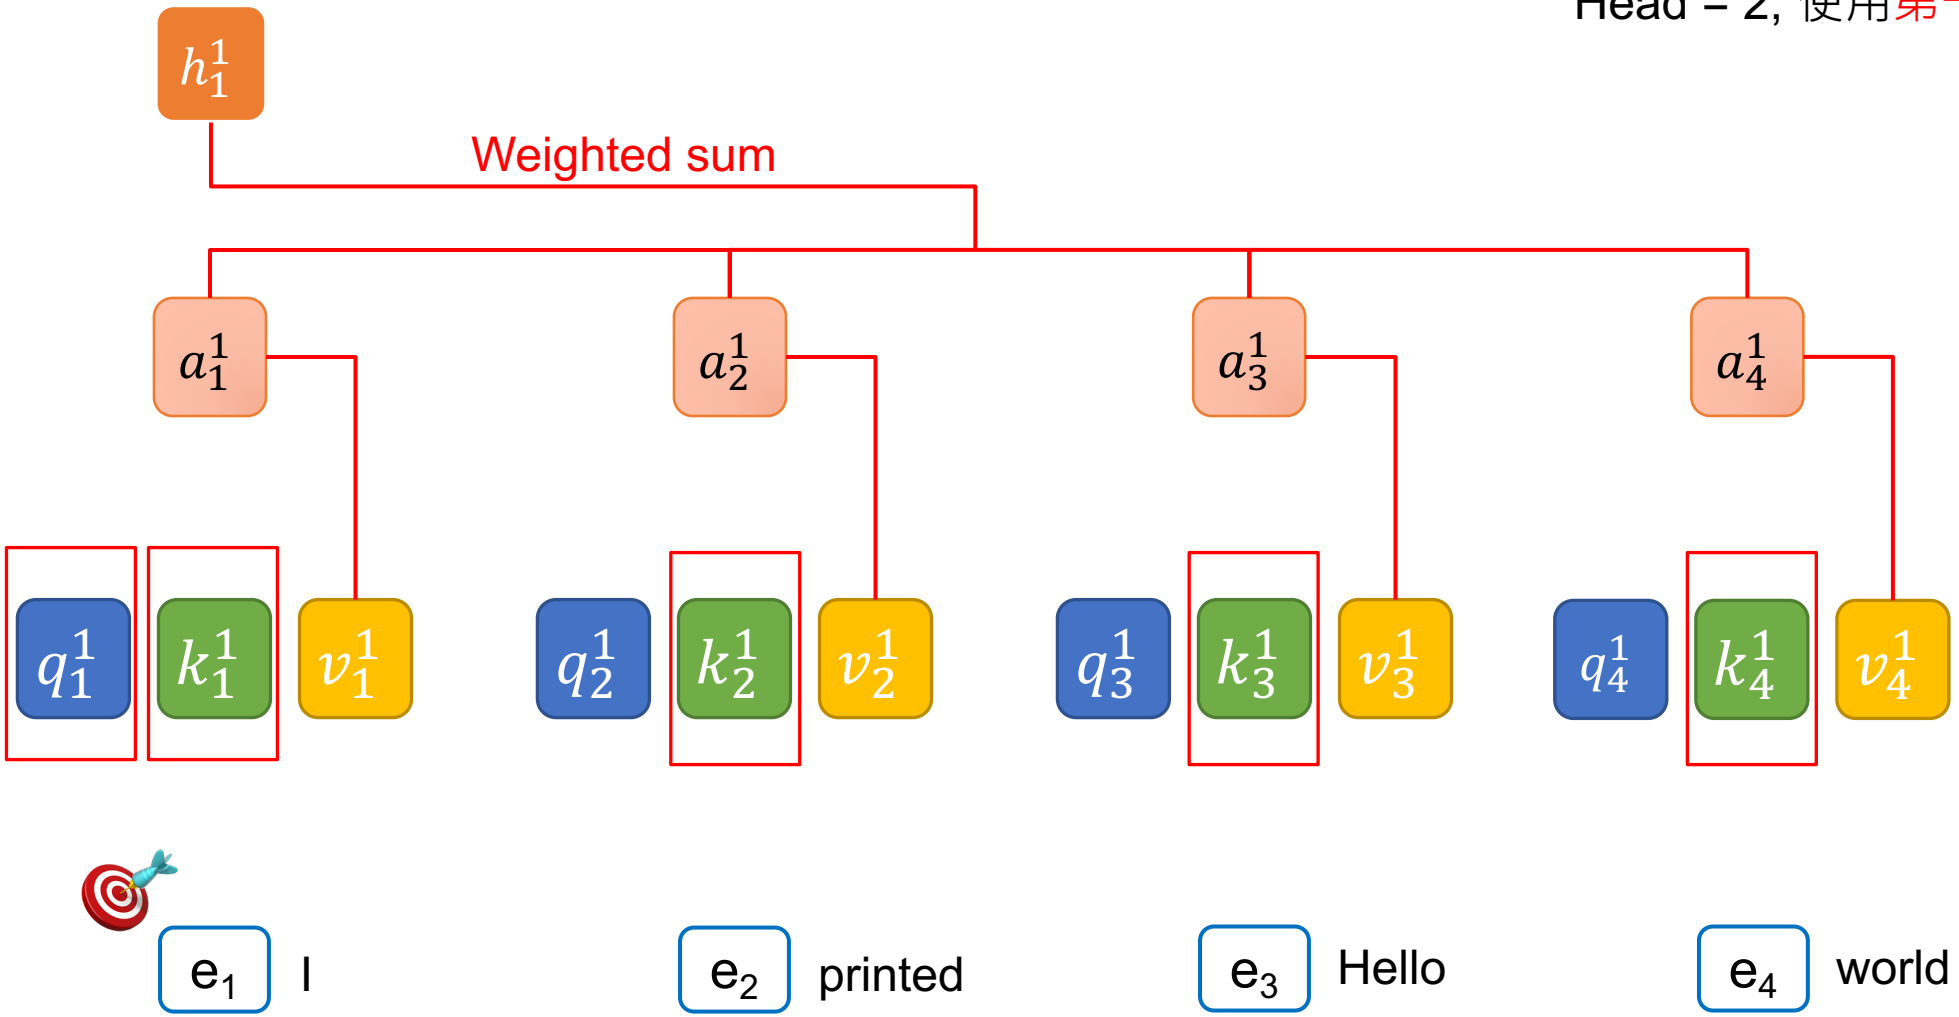

Query (Q), Key (K), and Value (V) with MHA

Head = 2, 使用第二個 head

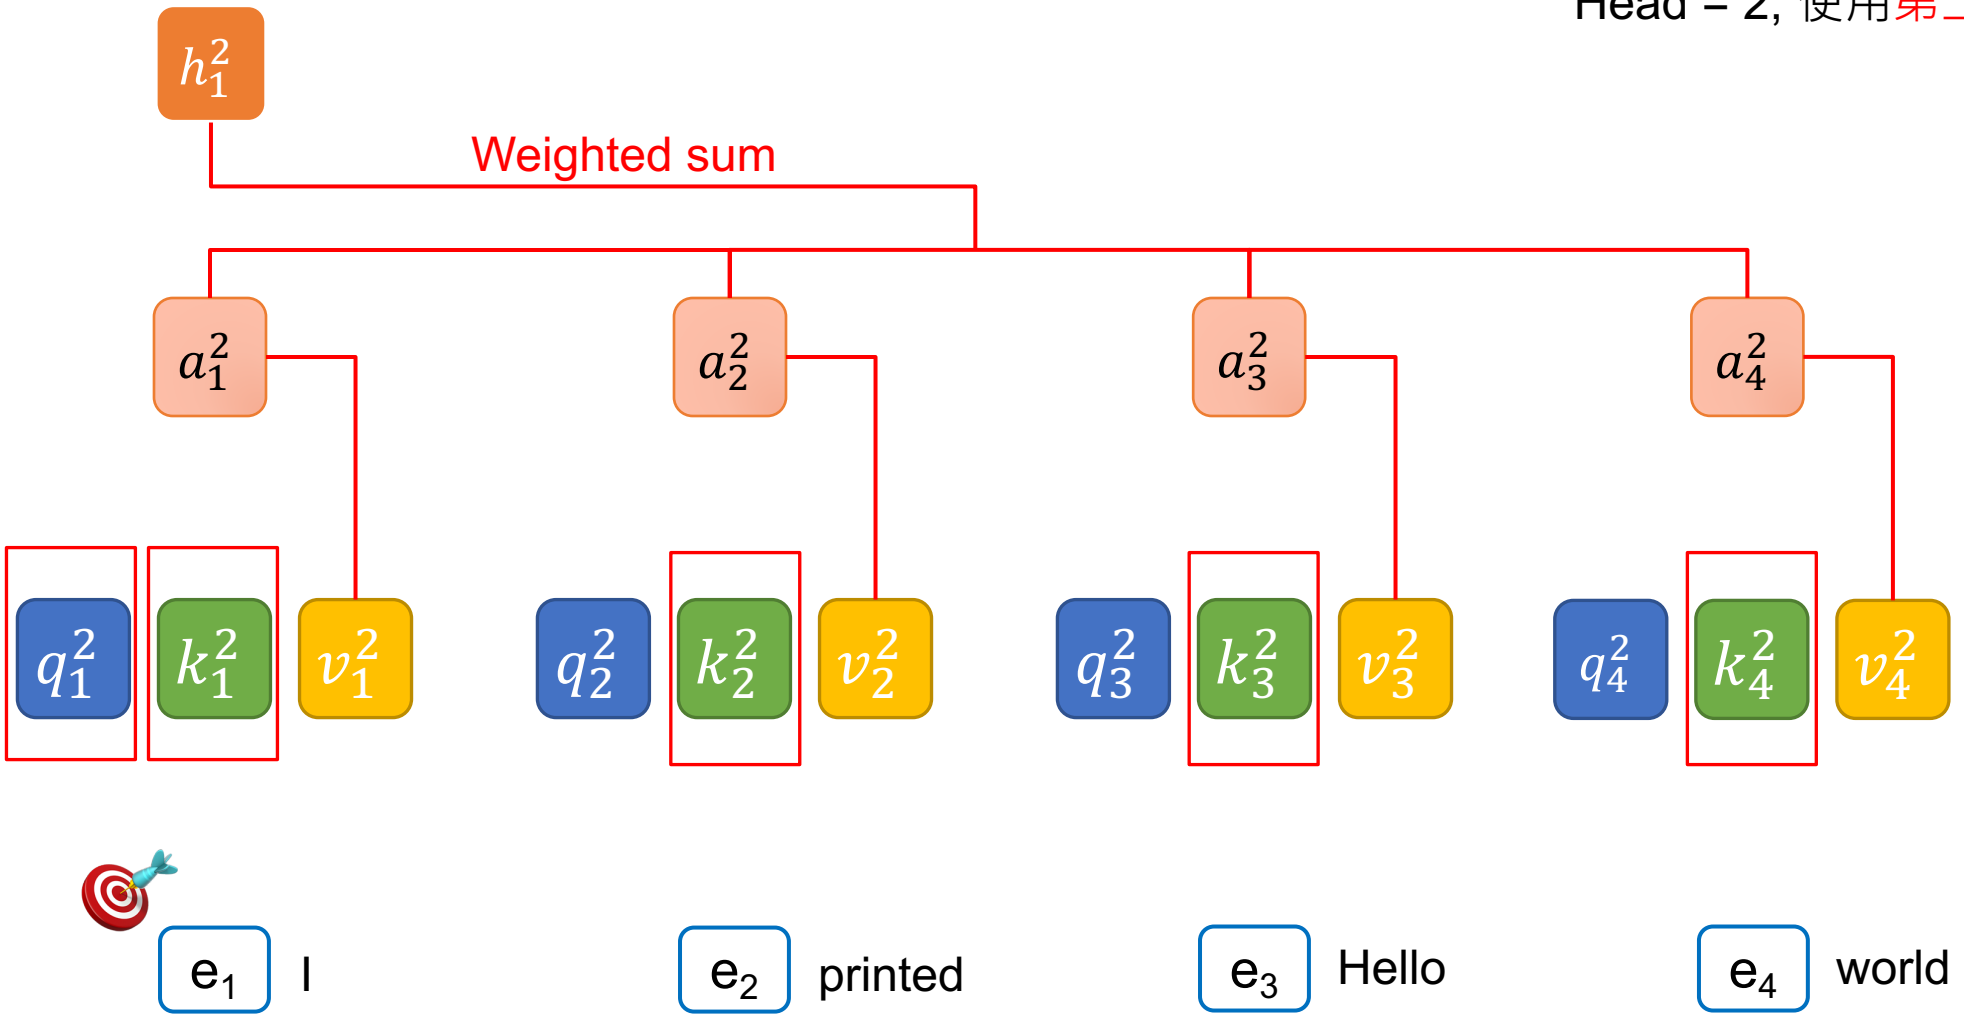

Input Dimensions of MHA

Head = 1 (without MHA)

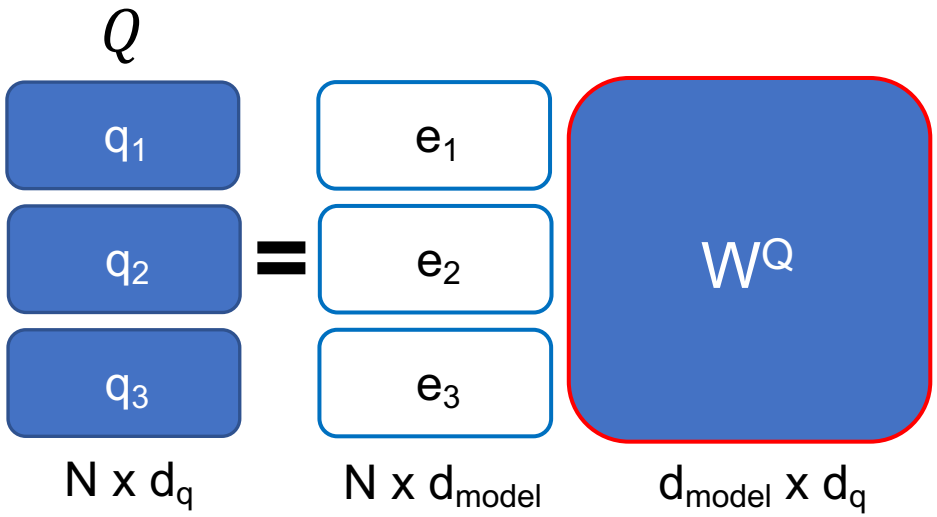

Head = 2 (with MHA)

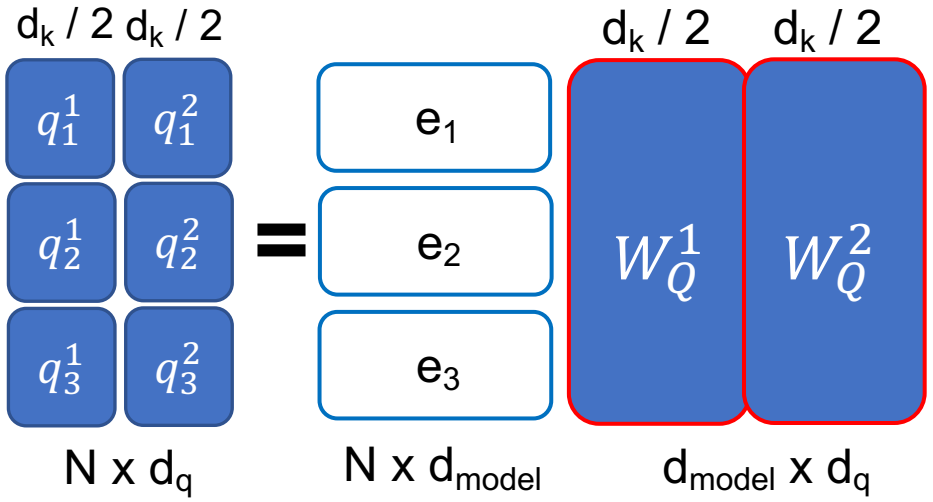

Output Dimensions

$$\text{Multihead}(Q, K, V) = \underbrace{\text{Concat}(\text{head}_1, \ldots, \text{head}_h)}_{\in\ \mathbb{R}^{N \times (h \cdot d)}} \underbrace{W^O}_{\in\ \mathbb{R}^{(h \cdot d) \times d_{out}}}$$

Head = 1 (without MHA)

Head = 2 (with MHA)

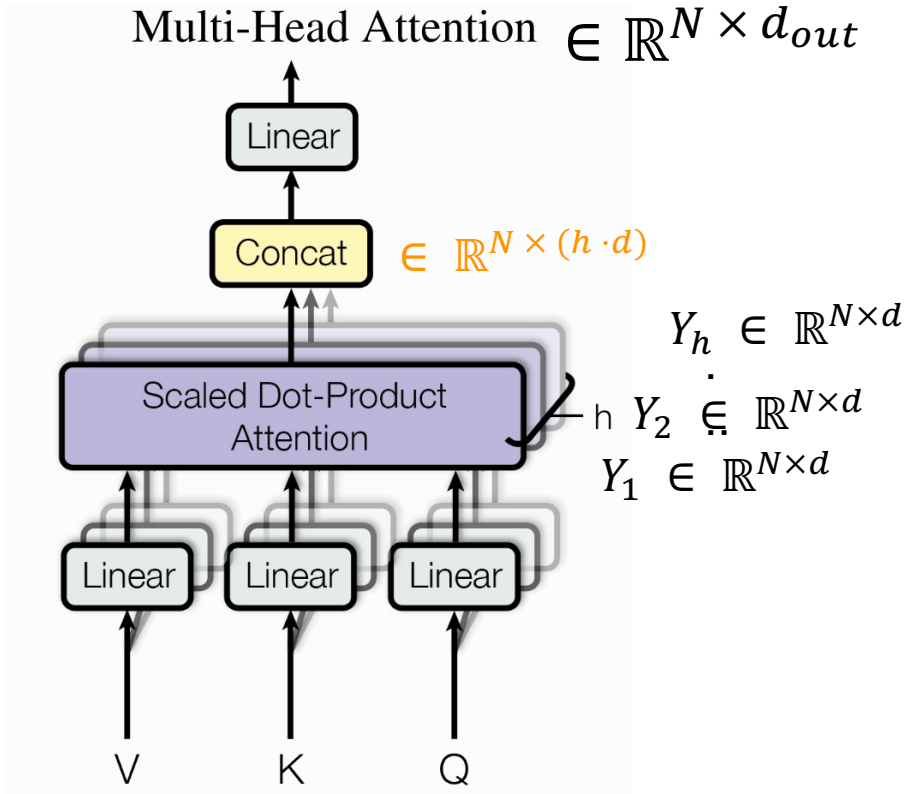

MHA 比較好嗎？

Clark, Kevin, et al. "What Does BERT Look at? An Analysis of BERT's Attention." Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP. 2019.

實驗證明：不同 head 關注的地方不太一樣

<layer>-<head number>

Head 8-10 關注受詞

- **Direct objects** attend to their verbs
- 86.8% accuracy at the `dobj` relation

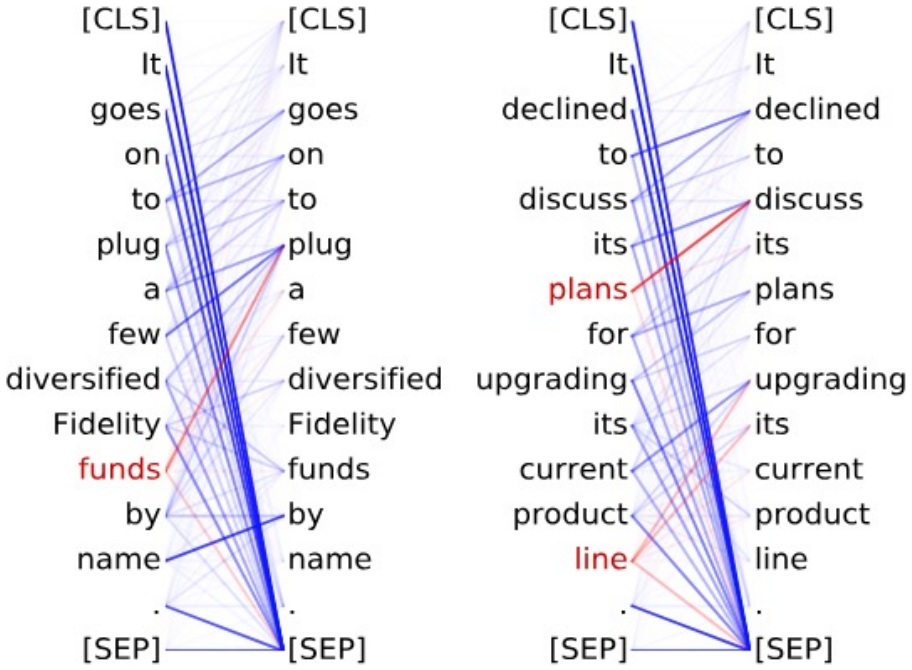

Head 8-11 關注限定詞

- **Noun modifiers** (e.g., determiners) attend to their noun
- 94.3% accuracy at the `det` relation

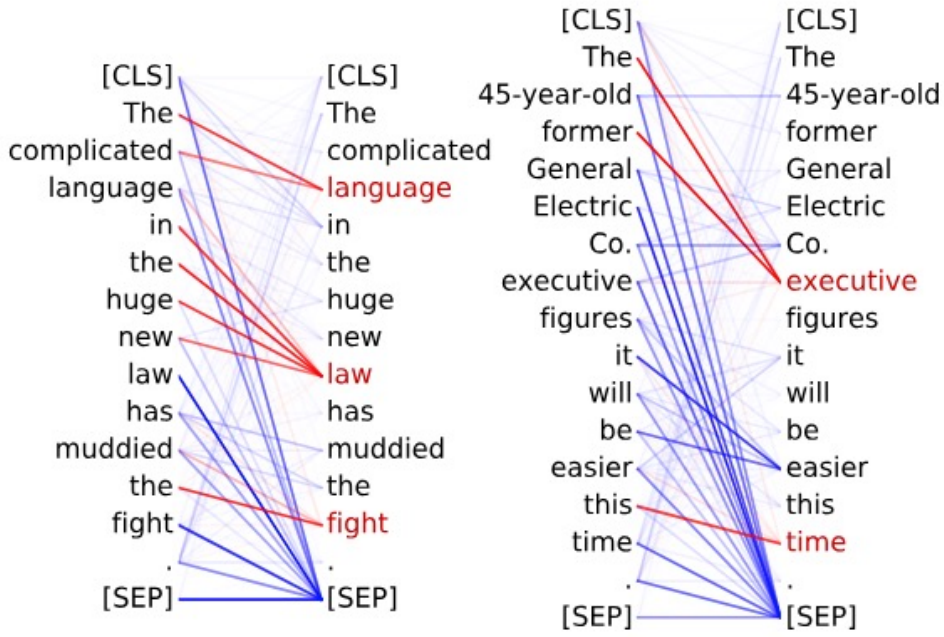

MHA 的過程中涵蓋了多次 softmax

每次 scaled dot-product attention 都會做 softmax

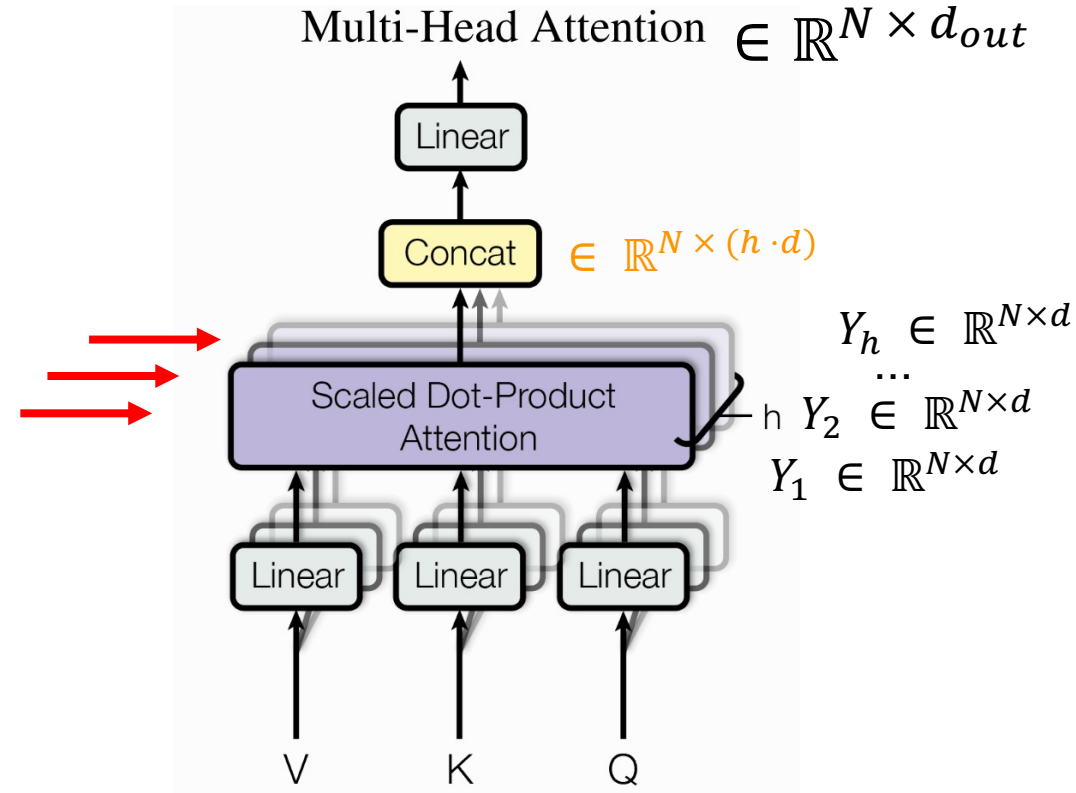

Summary of MHA

- 透過 MHA，模型會將表示投影到多個不同的子空間，讓每個 head 各自進行 self-attention，因此同一個時間點可以同時從不同角度捕捉序列中的關係
- 標準 Transformer 的 num_heads 通常不改變最終輸出維度，因為總模型維度 d_model 固定，head 數改變時會相應調整每個 head 的維度

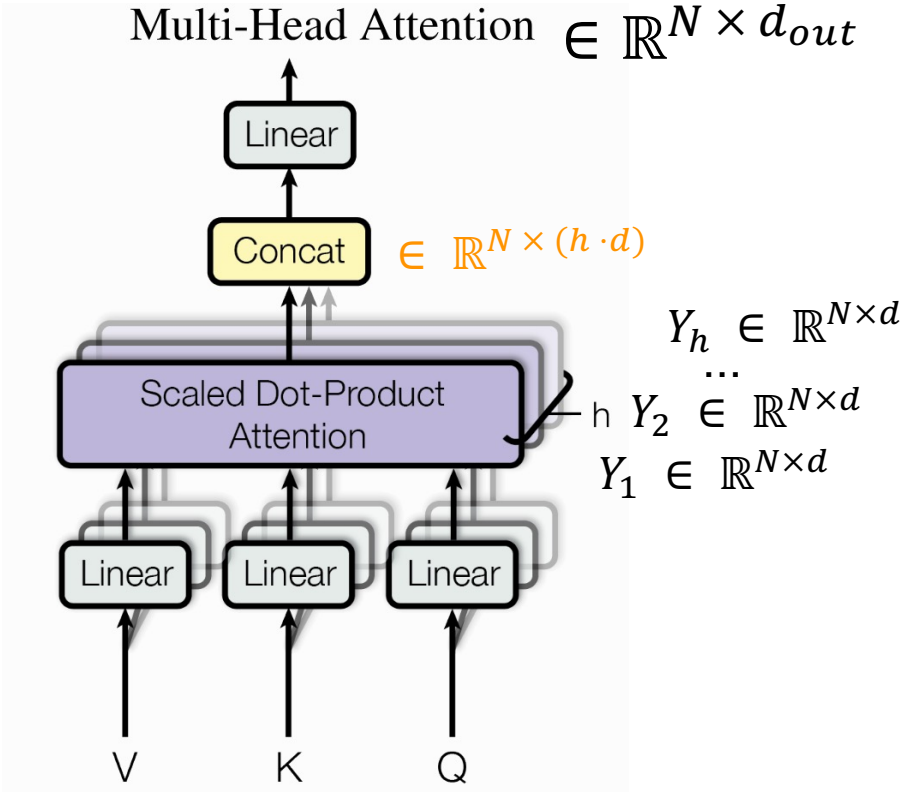

Outline of Transformer

- Add: Residual Connection
- Norm: Layer Normalization

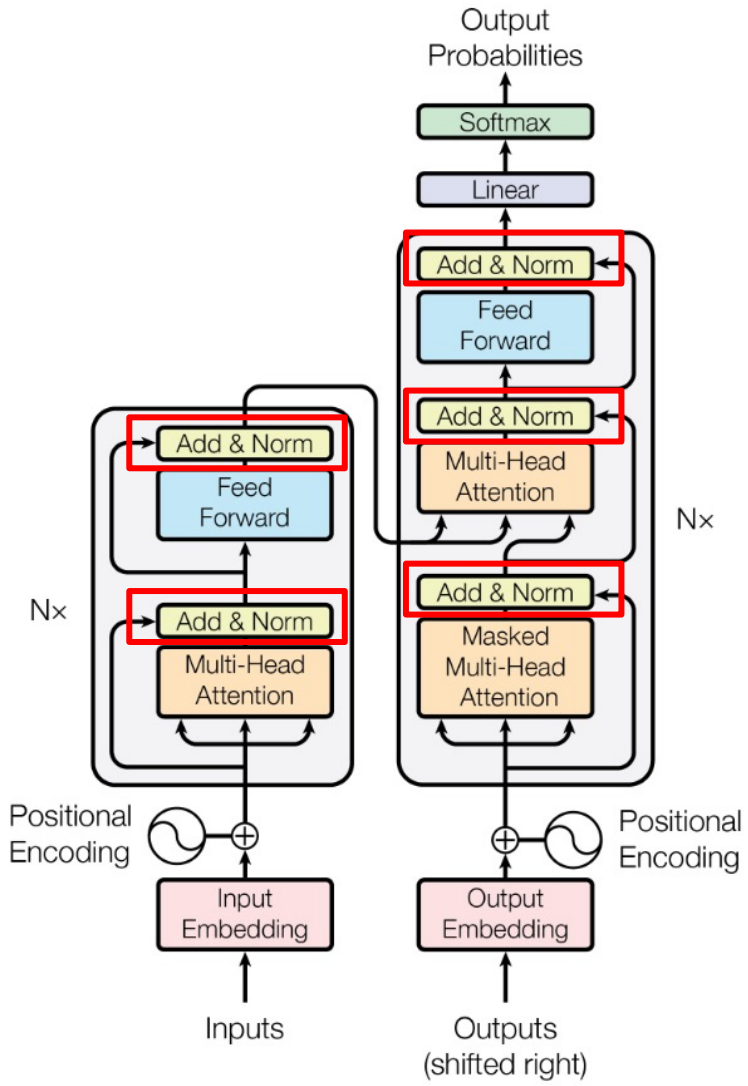

Vaswani, Ashish, et al. "Attention is all you need." NeurIPS (2017).

Add & Norm

- Research shows that Residual Connection (He, Kaiming, et al., 2016) stabilizes training.
- Let the layers in between learn the residual relationship (i.e. $Y - X$).

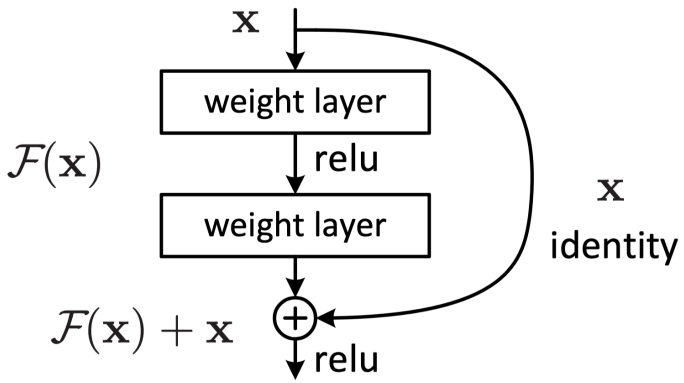

Figure 2. Residual learning: a building block.

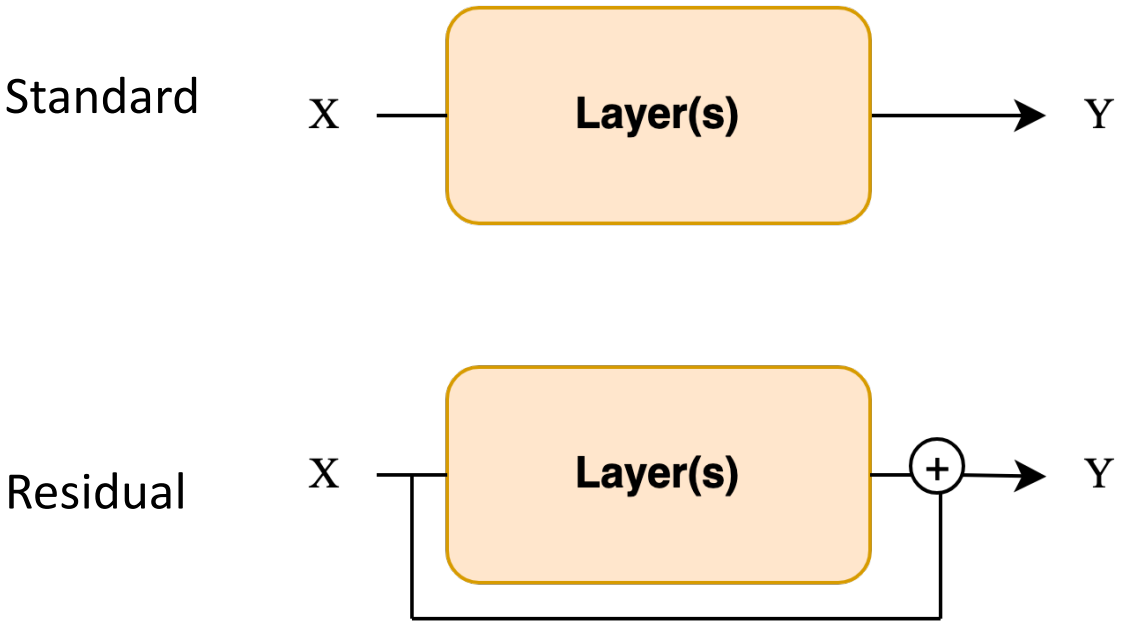

He, Kaiming, et al. "Deep residual learning for image recognition." Proceedings of the IEEE conference on computer vision and pattern recognition (2016).

Add & Norm (Layer Normalization)

- Layer Norm: For an input vector, calculate its mean and std across embedding dimension (within an example).
  - $\mathbf{norm}(X) = \frac{X-\mu}{\sigma}$

B: Batch Size
C: Channels / Feature Dimension
N: Sequence length (CNN 中代表長x寬的數目)

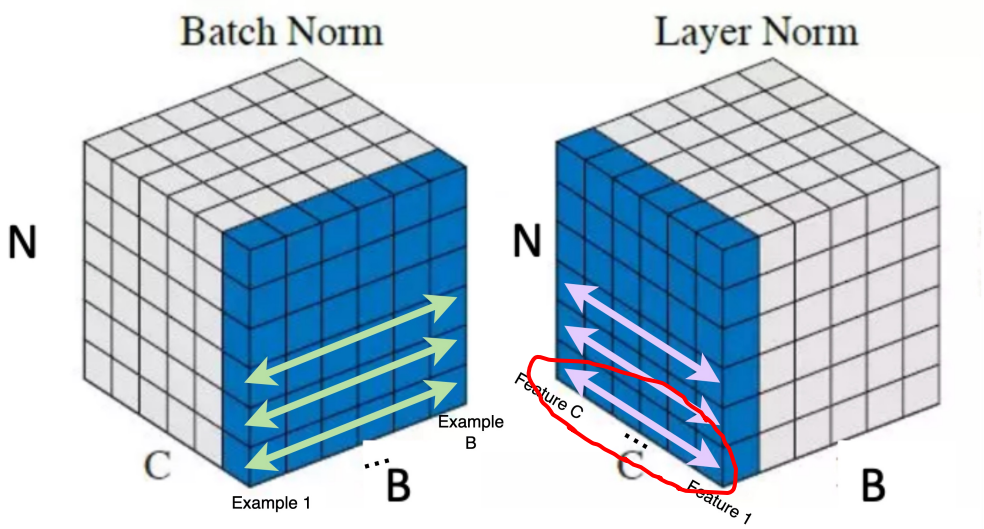

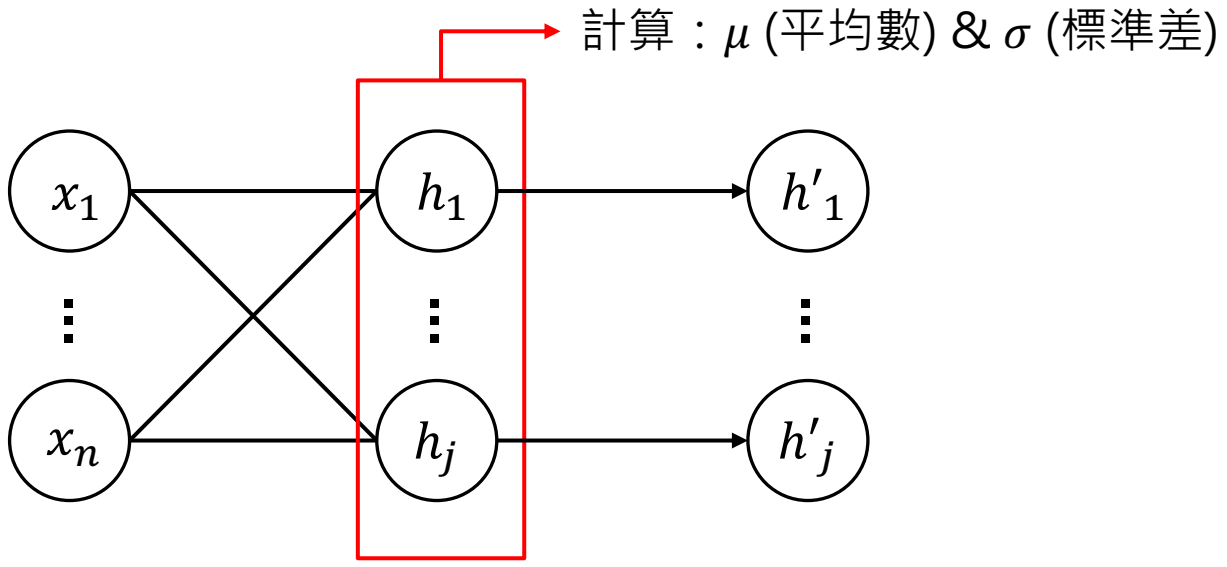

(Batch Normalization)

Ba, Jimmy Lei, Jamie Ryan Kiros, and Geoffrey E. Hinton. "Layer normalization." arXiv preprint arXiv:1607.06450 (2016).

Add & Norm (Layer Normalization)

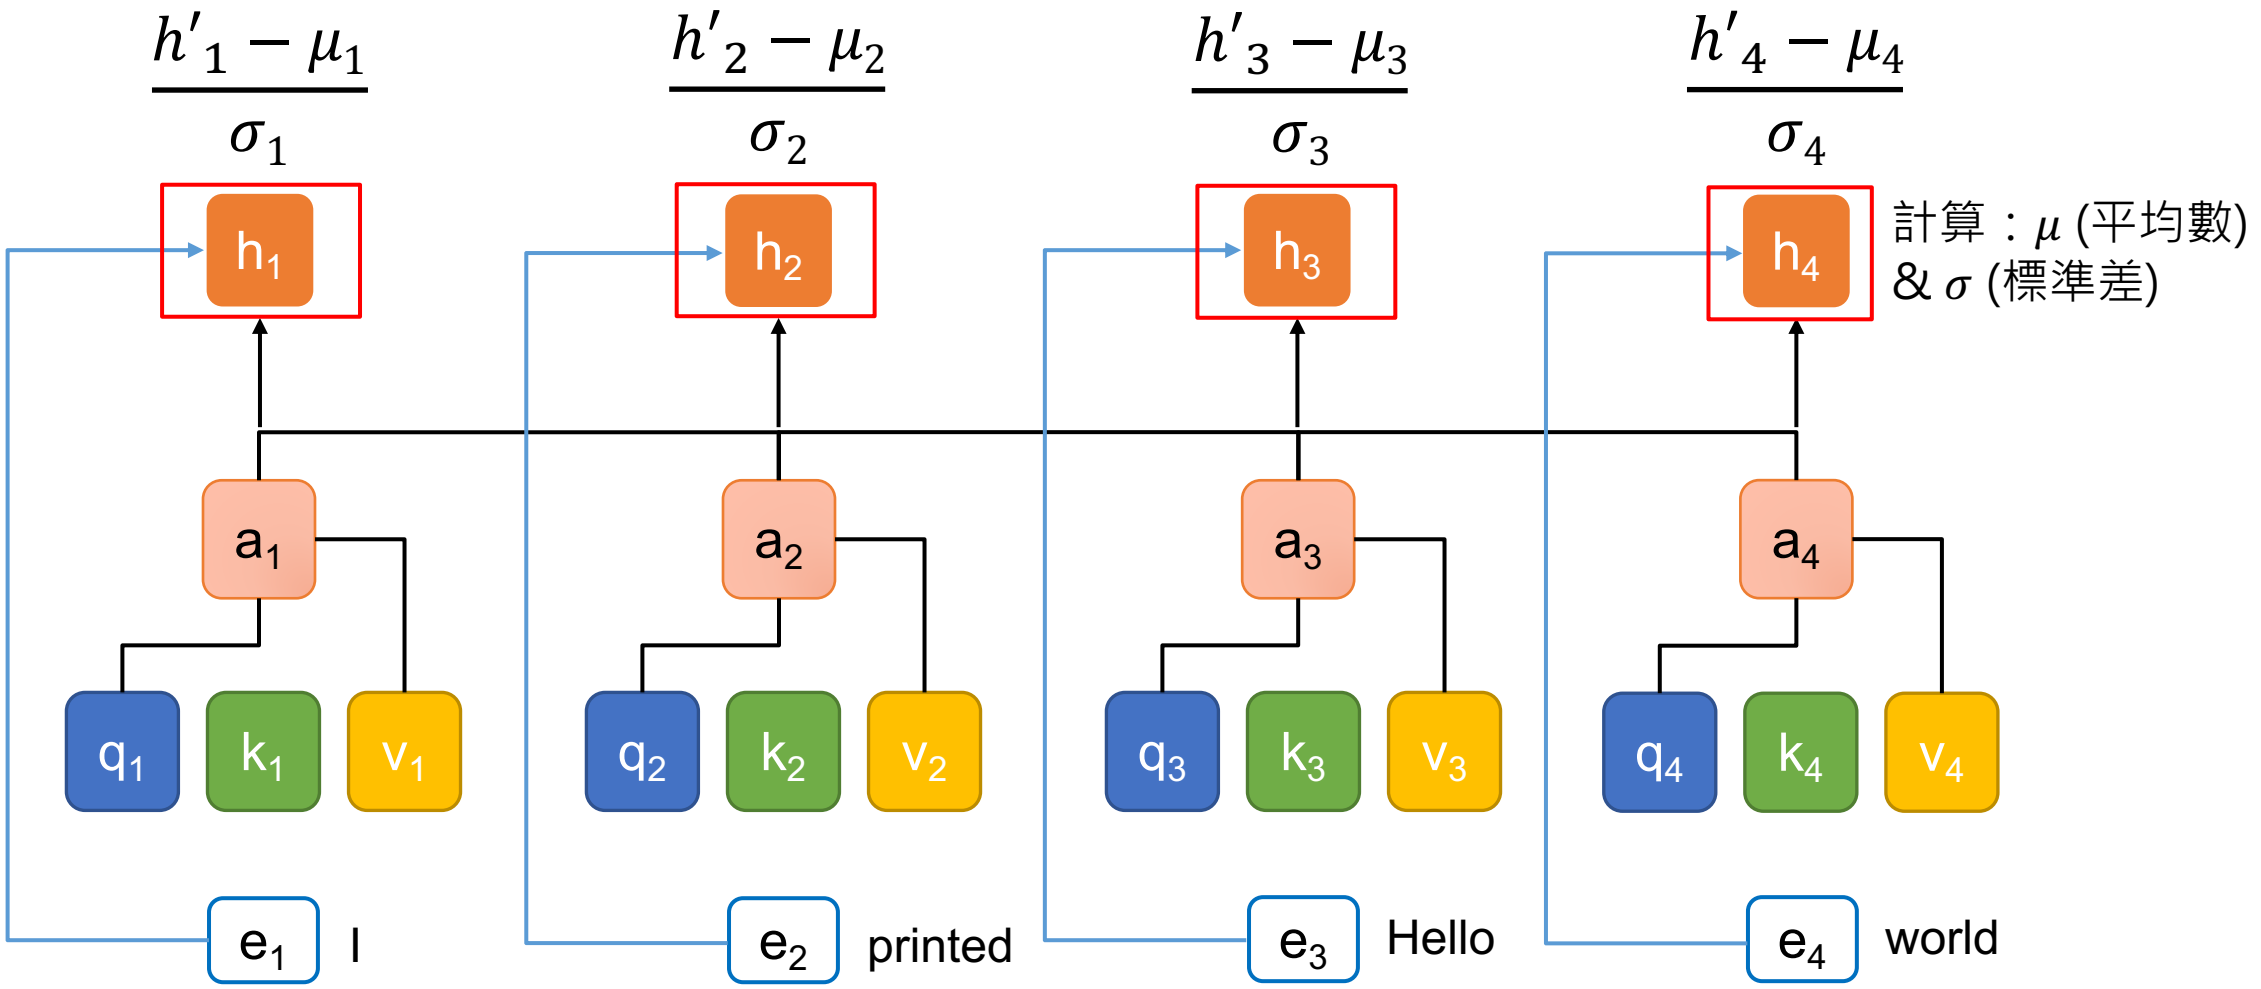

Ba, Jimmy Lei, Jamie Ryan Kiros, and Geoffrey E. Hinton. "Layer normalization." arXiv preprint arXiv:1607.06450 (2016).

Why LN not BN?

- 資料長度差距
- 模型實作表現差異

[Intro] Static padding

batch 1

I	love	CGU	[PAD]	[PAD]	[PAD]	
Yesterday	I	watched	a	good	movie	longest
Your	new	shoes	look	nice	[PAD]	

batch 2

Merry	Christmas	[PAD]	[PAD]	[PAD]	[PAD]
I	love	coding	[PAD]	[PAD]	[PAD]
Me	[PAD]	[PAD]	[PAD]	[PAD]	[PAD]

Outline of Transformer

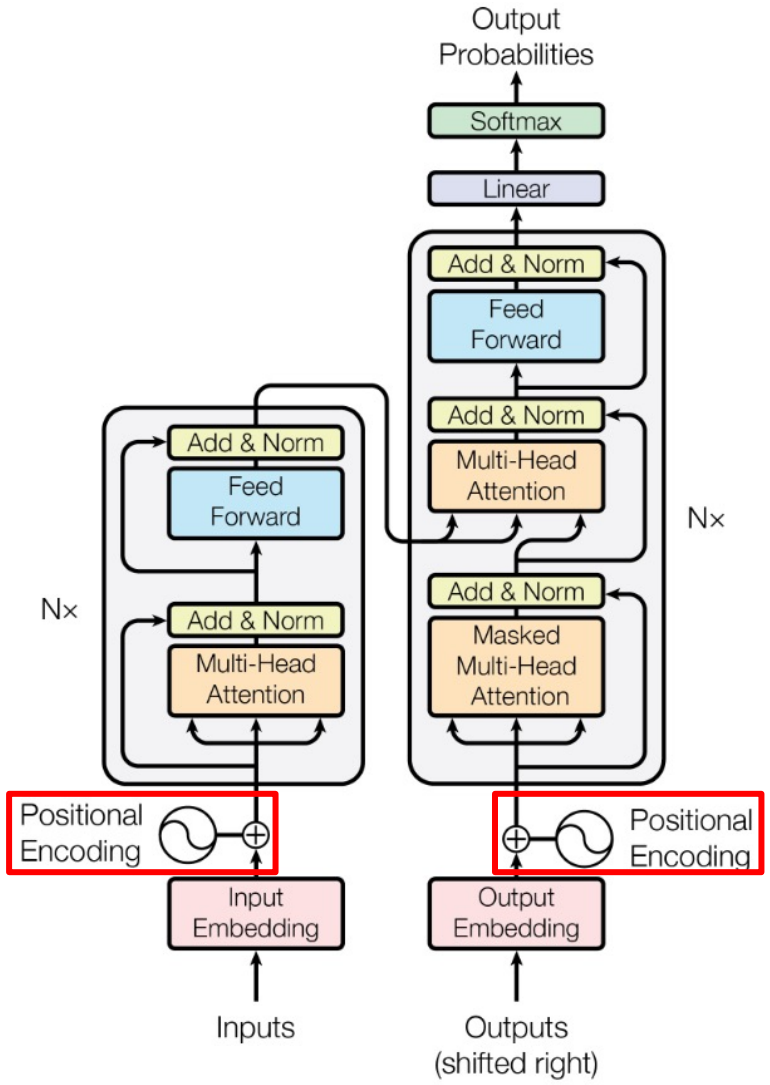

Vaswani, Ashish, et al. "Attention is all you need." NeurIPS (2017).

Lack of Position Modeling

- RNN 有位置資訊 (因為 RNN step-by-step 處理文字)
- 然而，Transformer 中 self-attention 的操作遺失了 step-by-step 處理文字的方式，故缺少位置資訊

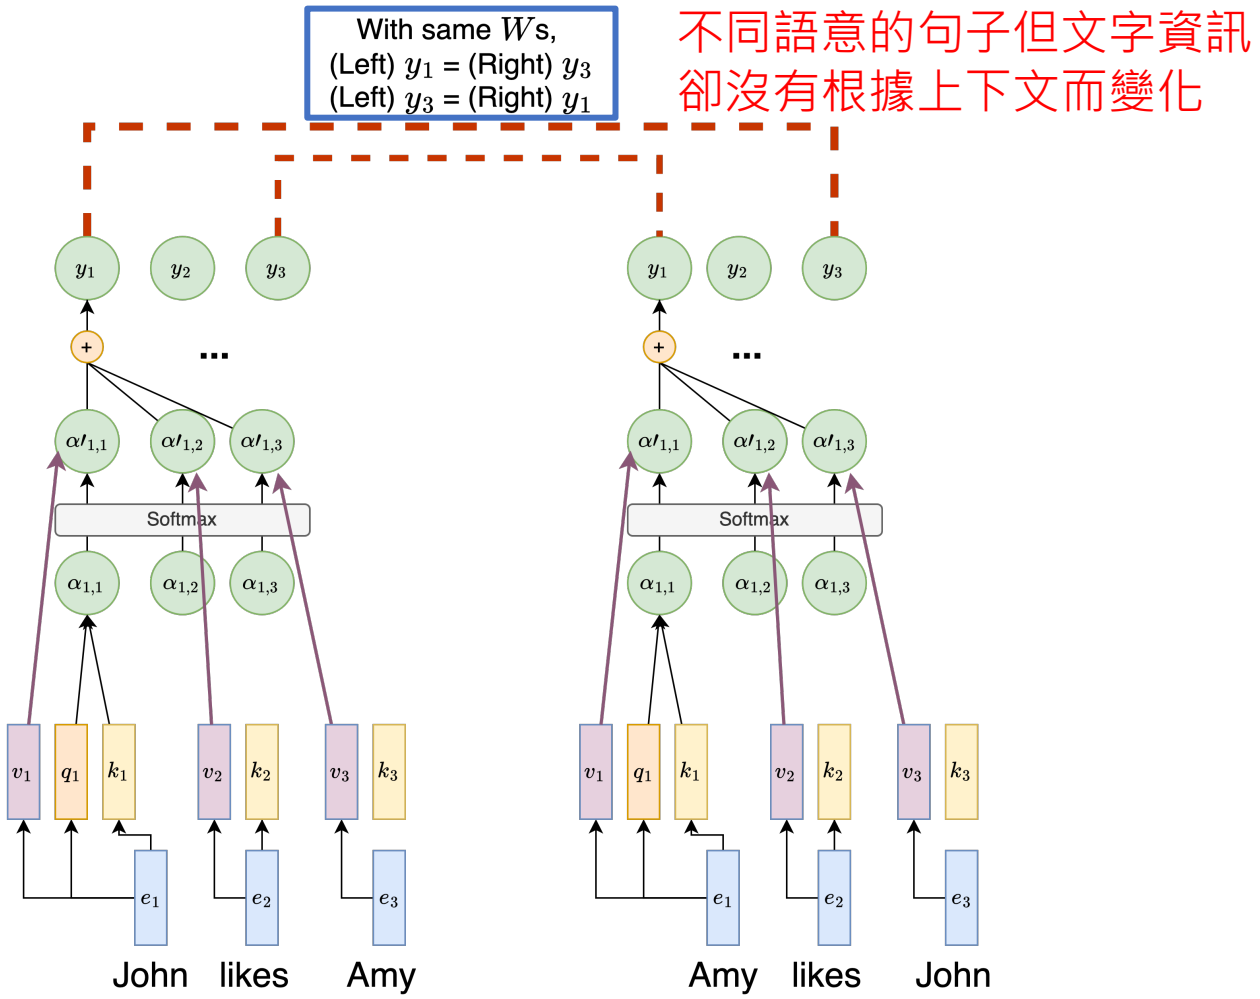

常見位置資訊作法：Learned Positional Embeddings

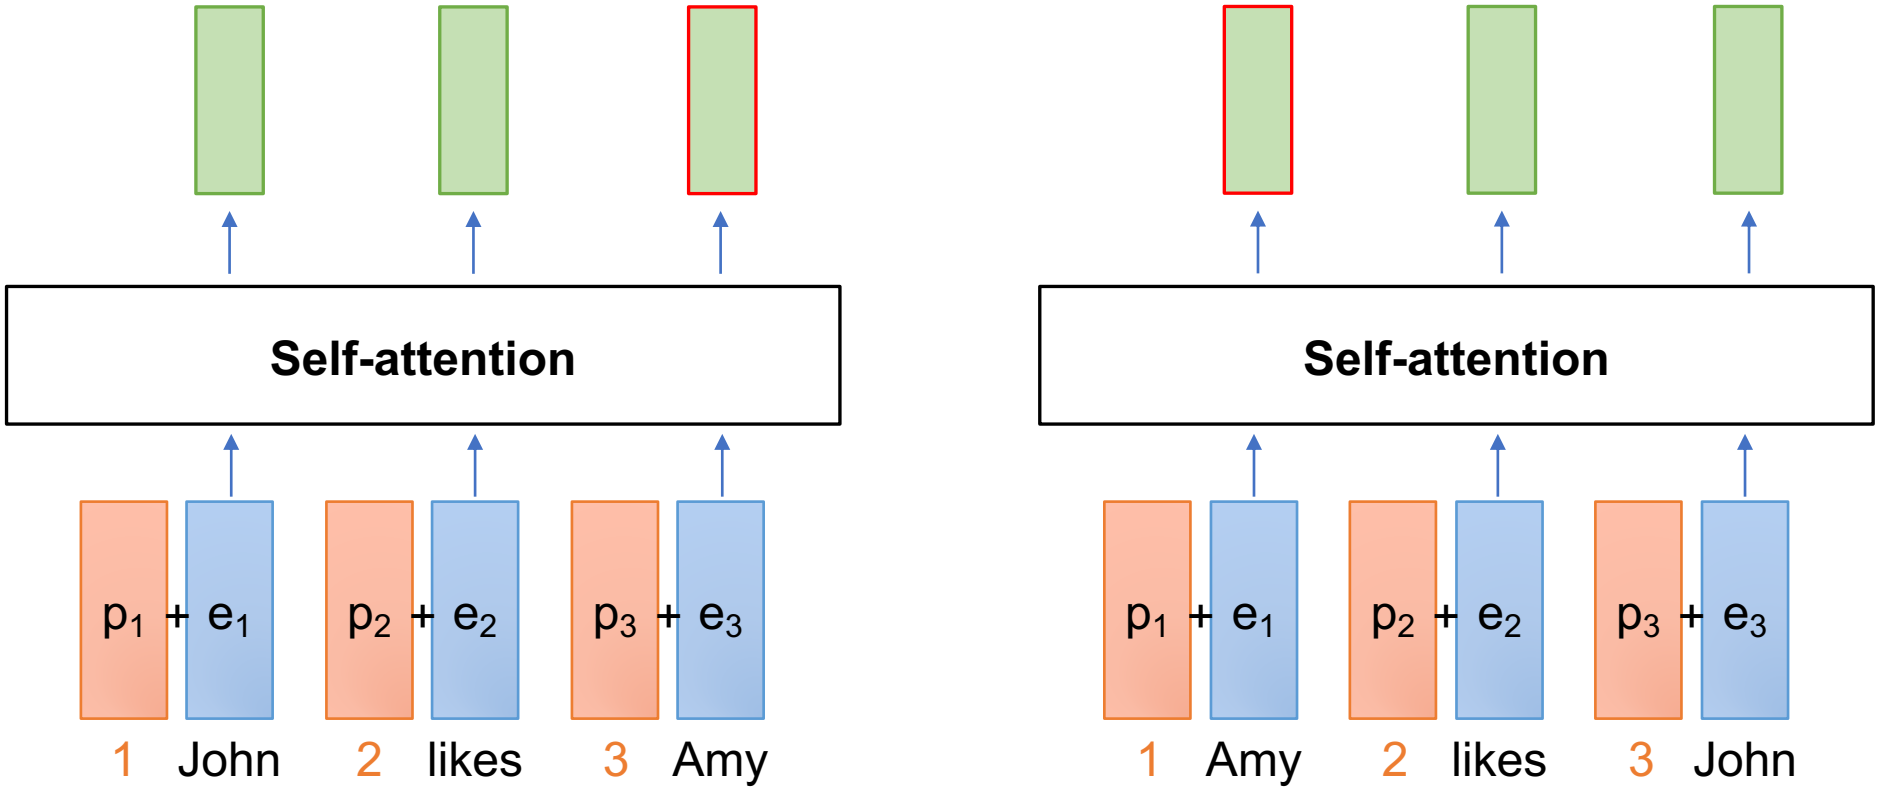

此做法又稱絕對位置資訊編碼 (Absolute positional Encoding)

Learned Positional Embeddings 程式碼

https://huggingface.co/transformers/v2.1.1/_modules/transformers/modeling_bert.html

```
class BertEmbeddings(nn.Module):
    """Construct the embeddings from word, position and token_type embeddings.
    """
    def __init__(self, config):
        super(BertEmbeddings, self).__init__()
        self.word_embeddings = nn.Embedding(config.vocab_size, config.hidden_size, padding_idx=0)
        self.position_embeddings = nn.Embedding(config.max_position_embeddings, config.hidden_size)
        self.token_type_embeddings = nn.Embedding(config.type_vocab_size, config.hidden_size)

        # self.LayerNorm is not snake-cased to stick with TensorFlow model variable name and be able to load
        # any TensorFlow checkpoint file
        self.LayerNorm = BertLayerNorm(config.hidden_size, eps=config.layer_norm_eps)
        self.dropout = nn.Dropout(config.hidden_dropout_prob)
```

缺點：Learned Positional Embeddings 不會超出 max_position 之後的資訊

Positional Encoding in Transformer

- 創造出相對位置的 embeddings
- 核心精神：利用三角函數的週期性來建構序列資訊
  - 偶數的 embedding 單元 (2i , i = 0, 2, 4, …, d_{model} / 2 - 2)：Sin 函數
  - 奇數的 embedding 單元 (2i+1, i = 1, 3, 5, …, d_{model} / 2 - 1)： Cosine 函數

$$PE_{(pos,2i)} = sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos,2i+1)} = cos(pos/10000^{2i/d_{\text{model}}})$$

i越大 -> 分母越大 -> 相位變化越慢 -> 頻率越小

10000是實驗得出的較佳數值

輸入序列中任一token的位置

Figure source: https://nlp.seas.harvard.edu/annotated-transformer/#positional-encoding

Positional Encoding Visualization

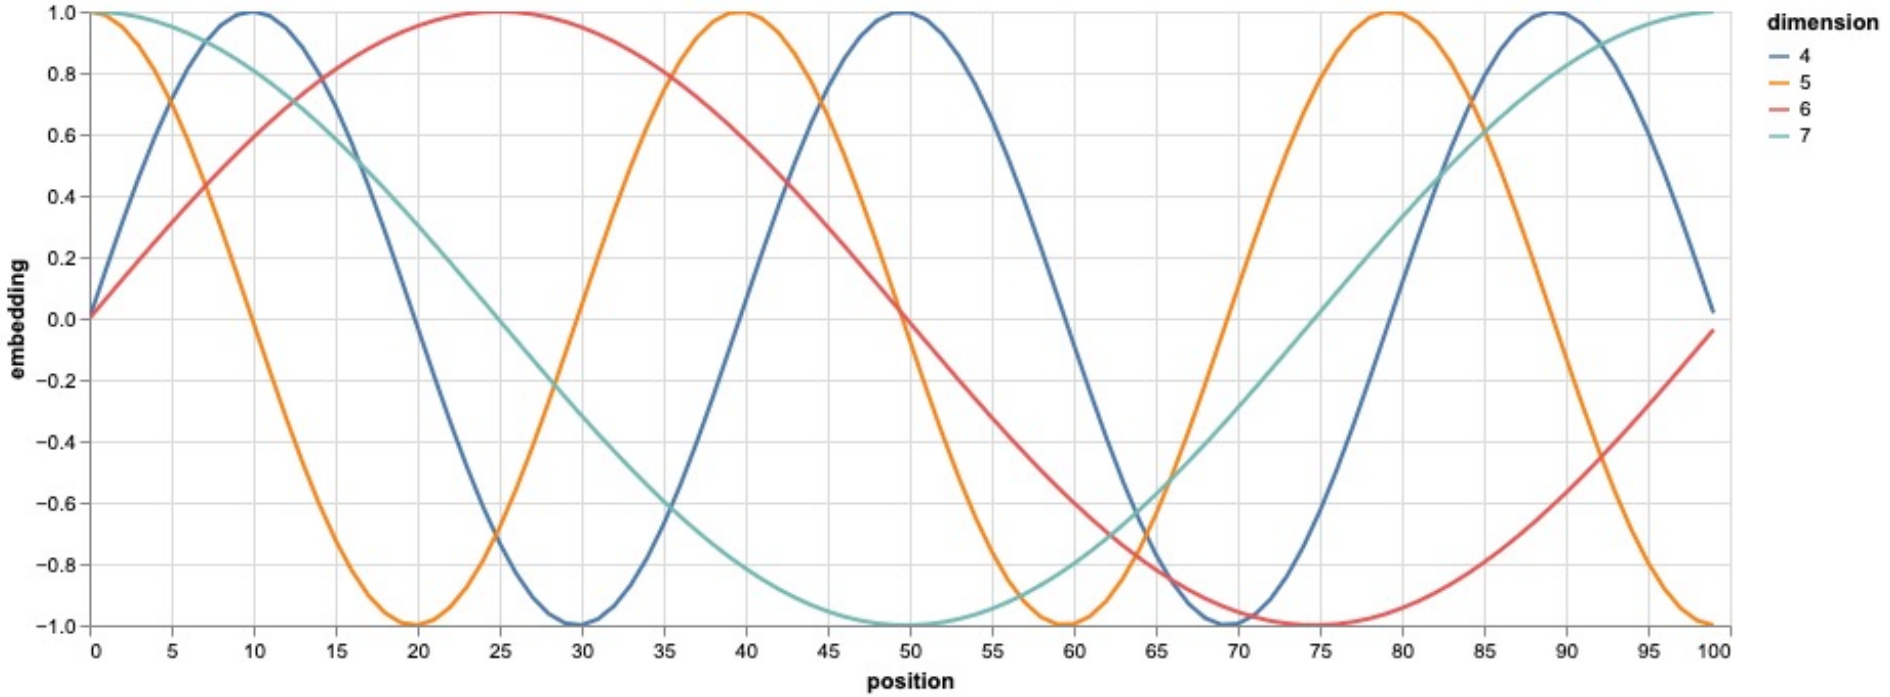

- Question: 為什麼需要同時有 Sin 和 Cosine 的函數？
- Ans: 藉由不同週期性來捕捉序列中各個位置的關係

$\sin 30^\circ = 0.5$
$\sin 150^\circ = 0.5$
vs.
$(\sin 30^\circ, \cos 30^\circ) = (0.5, 0.866)$
$(\sin 150^\circ, \cos 150^\circ) = (0.5, -0.866)$

Figure source: https://nlp.seas.harvard.edu/annotated-transformer/#positional-encoding

Encoder-Decoder 橋樑: Cross-Attention

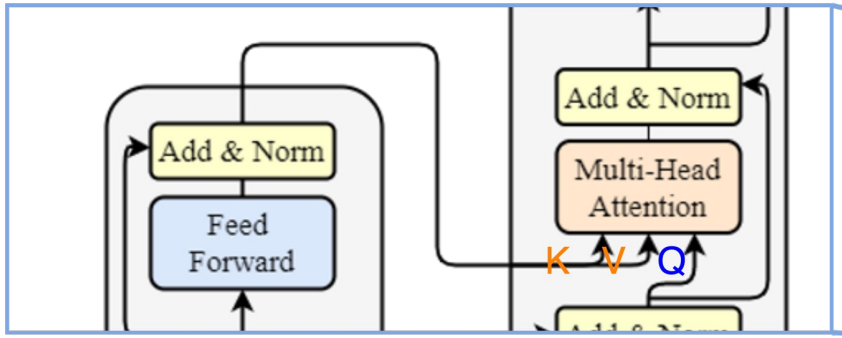

K, V uses the output from encoder as their X;

Q uses the decoder's information as X.

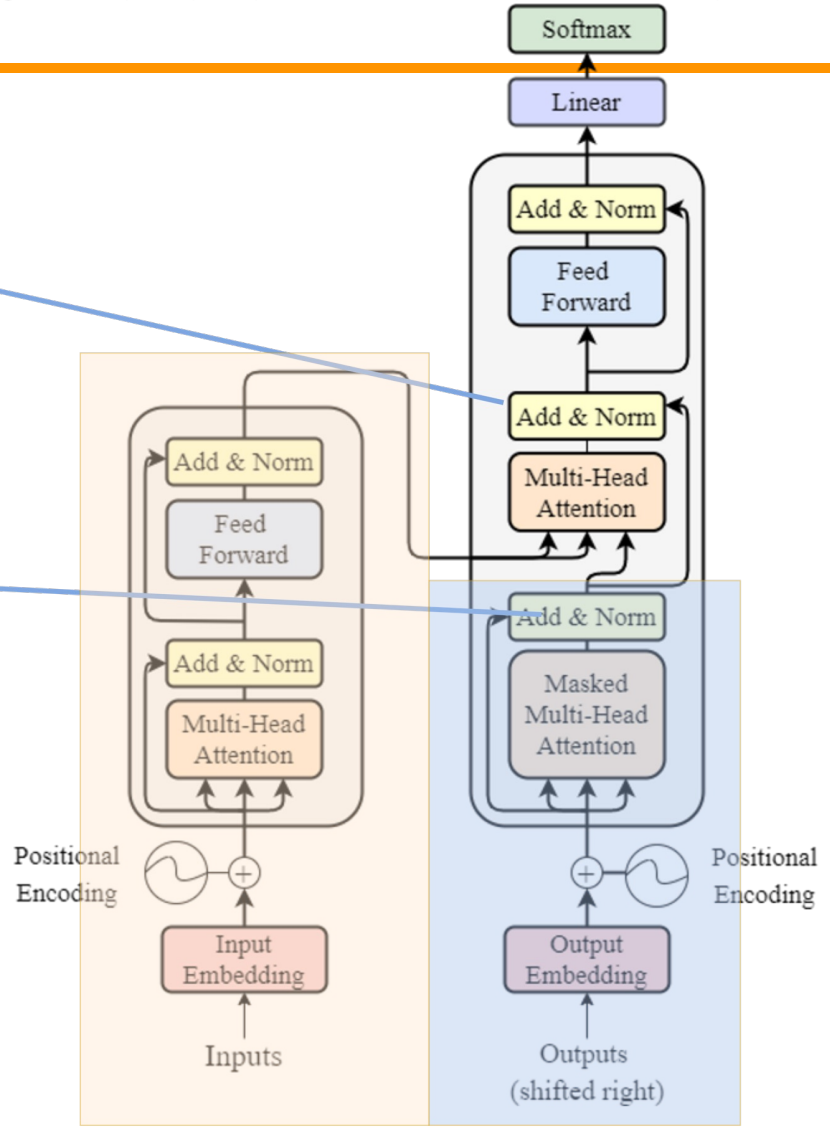

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

Cross-attention

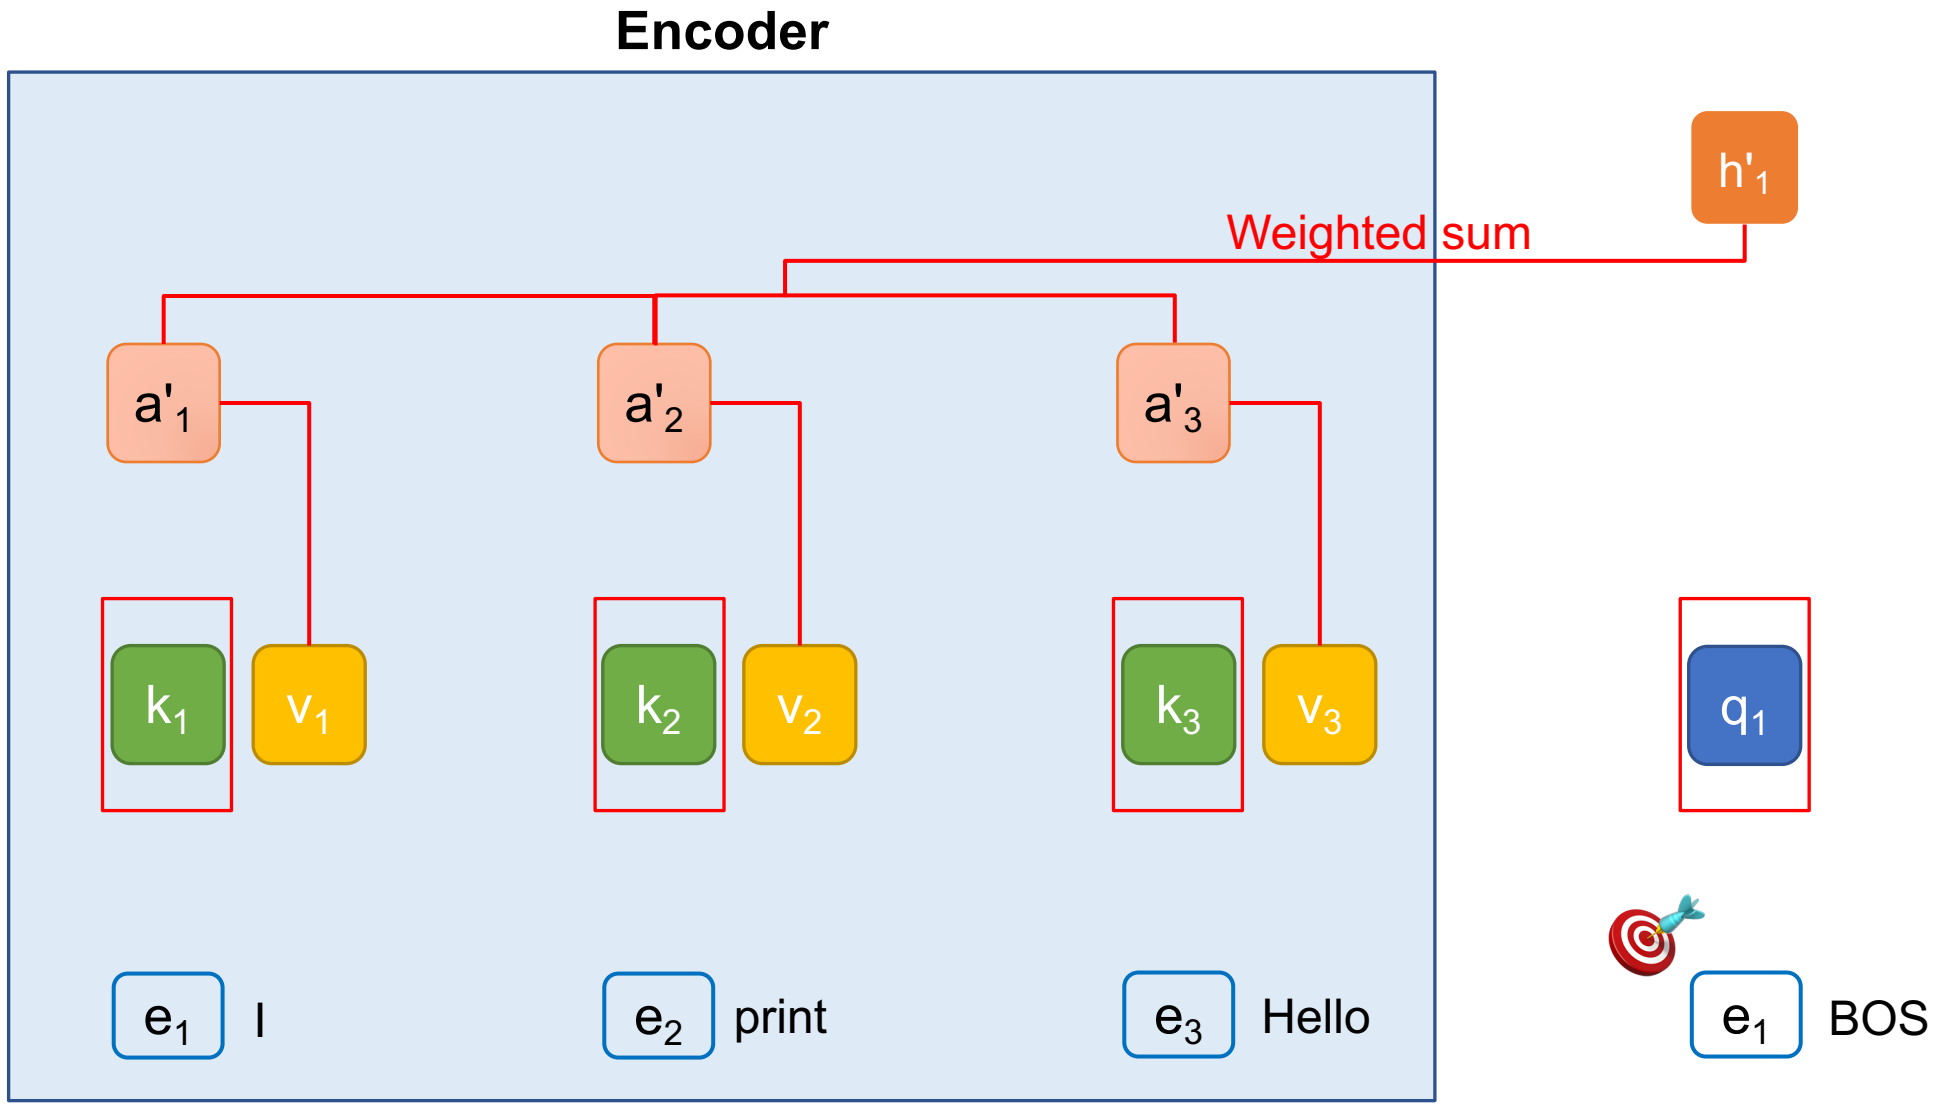

Number of Layers

Input, output dimension match!
We can have multiple encoders stacked up, to increase the number of trainable parameters.

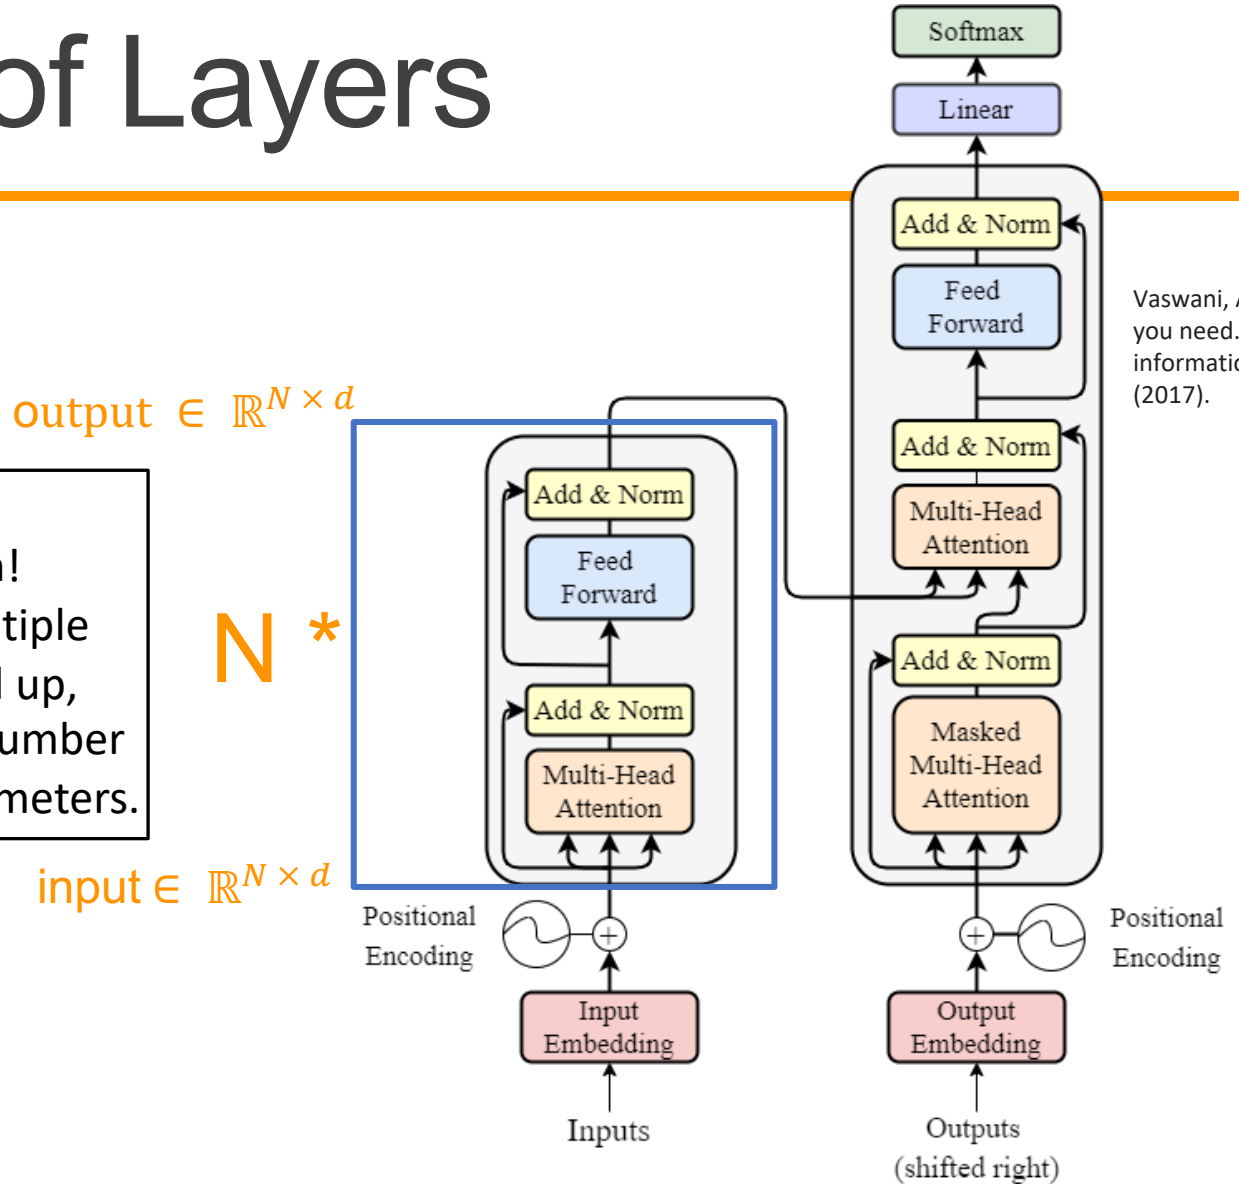

Vaswani, Ashish, et al. "Attention is all you need." Advances in neural information processing systems 30 (2017).

Outline of Transformer

Feed Forward Networks:

- Simply multiply by a weight matrix (MLP).
- Used to project to a specific dimension.

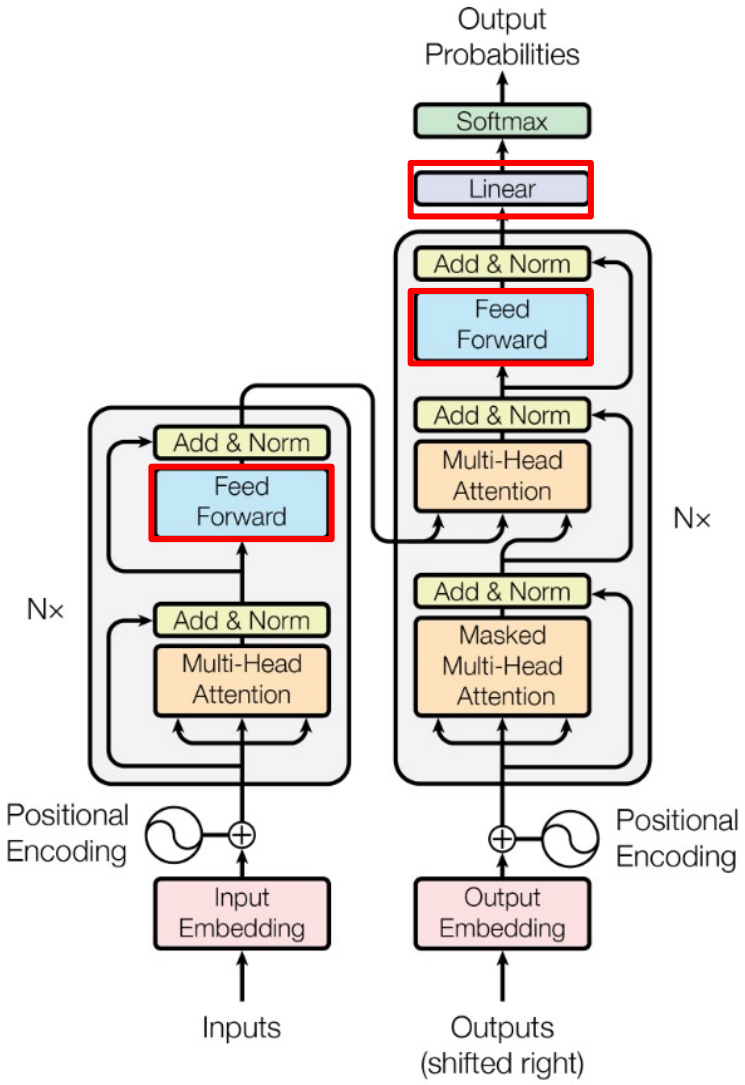

Vaswani, Ashish, et al. "Attention is all you need." NeurIPS (2017).

[Recap] Text Generation with RNN

*本頁的RNN不包含輸出層V

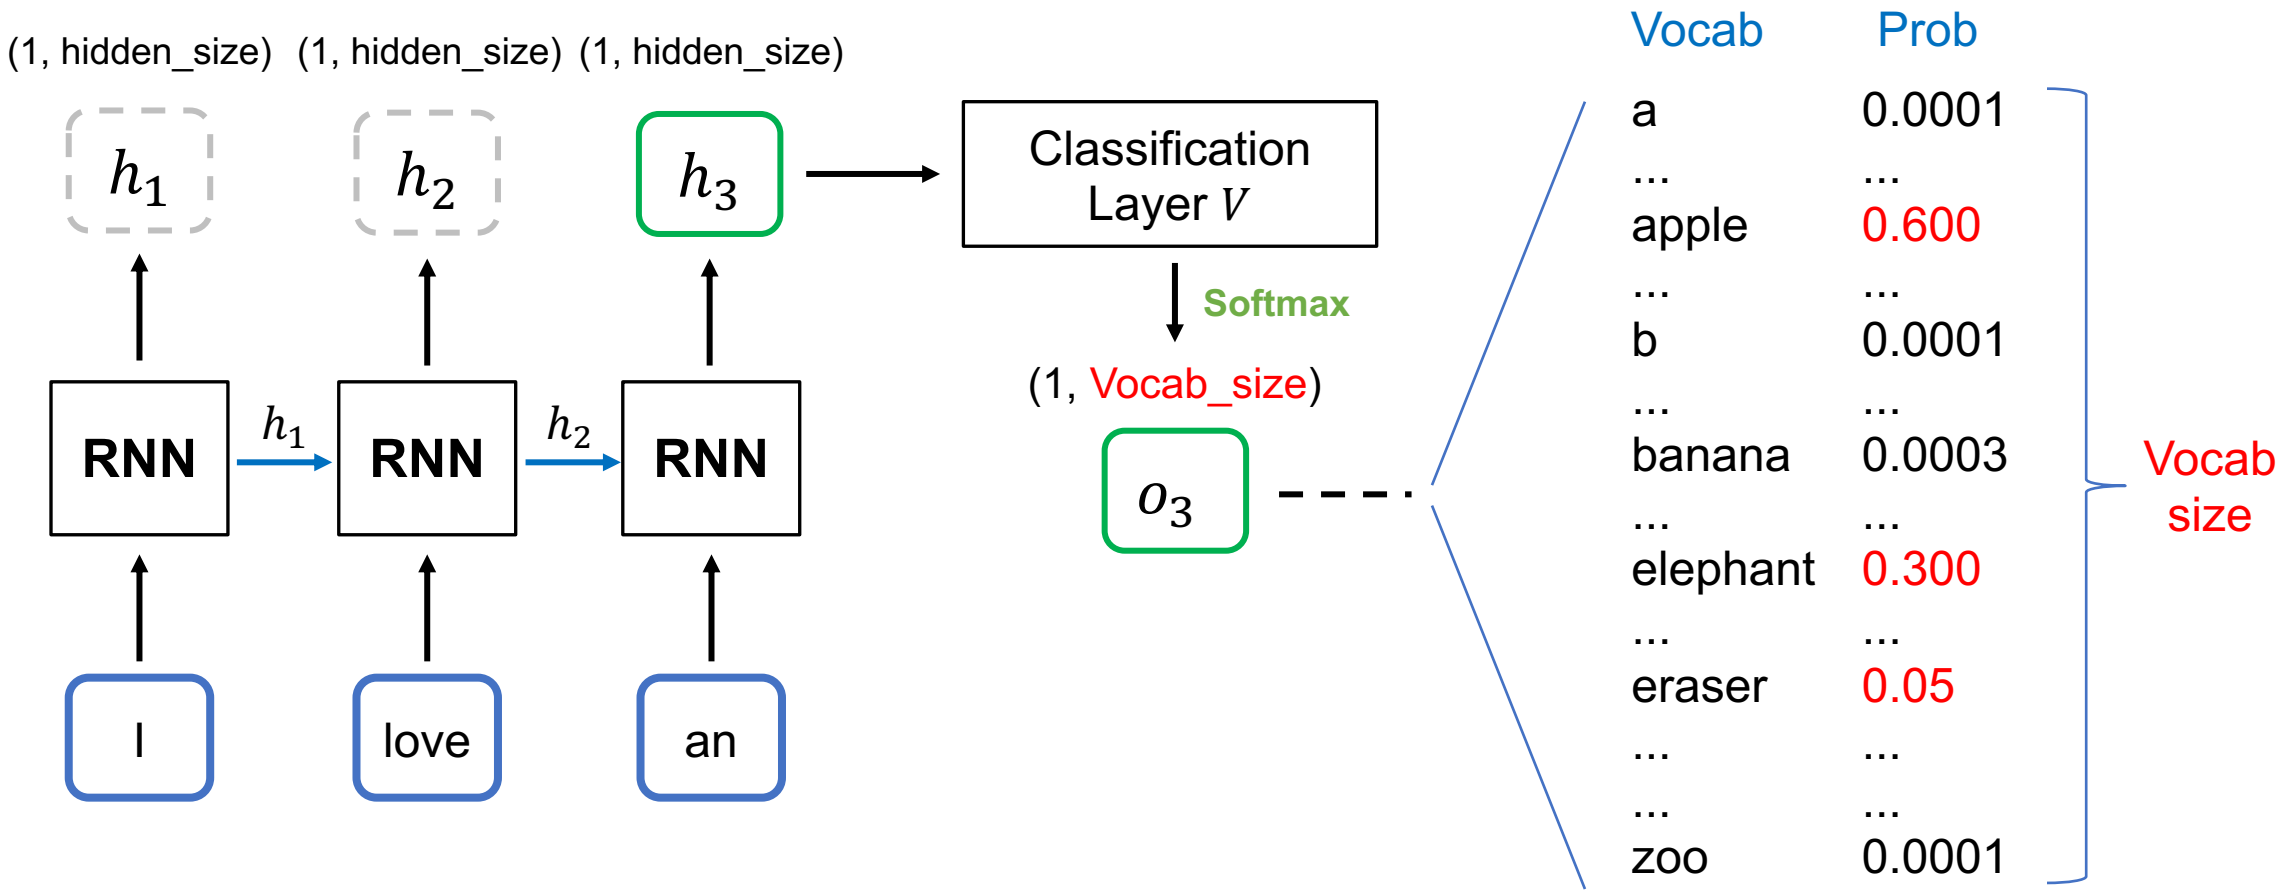

NLP

Word embeddings

Transformers' Achievements

1. In original paper, State-of-The-Art (SoTA) of machine translation.

Table 2: The Transformer achieves better BLEU scores than previous state-of-the-art models on the English-to-German and English-to-French newstest2014 tests at a fraction of the training cost.

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	**41.29**	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	**$3.3 \cdot 10^{18}$**	
Transformer (big)	**28.4**	**41.8**	$2.3 \cdot 10^{19}$	

6 encoders,
6 decoders,
h = 512, 100K train steps.

6 encoders,
6 decoders,
h = 1024, 300K train steps.

Vaswani, Ashish, et al. "Attention is all you need." Advances in neural information processing systems 30 (2017).

Transformer Variants (there are many more)

- Universal Transformer
  - Adds Adaptive Computation Time
  - Dehghani, Mostafa, et al. "Universal transformers." arXiv preprint arXiv:1807.03819 (2018).
- Longformer
  - Tackles long documents
  - Iz Beltagy, et al. (2020). Longformer: The Long-Document Transformer. arXiv:2004.05150.
- Roformer
  - Changes the design of positional encoding
  - Su, Jianlin, et al. "RoFormer: Enhanced Transformer with Rotary Position Embedding. CoRR abs/2104.09864 (2021)." arXiv preprint arXiv:2104.09864 (2021).
- Vision Transformer
  - Tackles vision tasks
  - Alexey Dosovitskiy, et al. “An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale.” International Conference on Learning Representations. 2021.

Additional Links

- https://nlp.seas.harvard.edu/annotated-transformer/
- https://datascience.stackexchange.com/questions/82451/why-is-10000-used-as-the-denominator-in-positional-encodings-in-the-transformer

Thank you!

Instructor: 林英嘉

✉ yjlin@cgu.edu.tw

TA: 林宣辰

✉ b1128015@cgu.edu.tw